

#Define Problem Statement and perform Exploratory Data Analysis

##Problem Statement:

Given a set of attributes for an Individual, determine if a credit line should be extended to them. If so, what should the repayment terms be in business recommendations?

Determine the creditworthiness of potential borrowers using various attributes, to ensure that the loans are given to those who are most likely to repay them.

From the company's perspective:

- LoanTap is at the forefront of offering tailored financial solutions to millennials.
- Their innovative approach seeks to harness data science for refining their credit underwriting process.
- The focus here is the Personal Loan segment. A deep dive into the dataset can reveal patterns in borrower behavior and creditworthiness.
- Analyzing this dataset can provide crucial insights into the financial behaviors, spending habits, and potential risk associated with each borrower.
- The insights gained can optimize loan disbursement, balancing customer outreach with risk management.

From the learner's perspective:

- Tackling this case offers practical exposure to real-world financial data and its challenges.
- Logistic Regression, a foundational algorithm, is pivotal in binary outcomes like loan decisions.
- Participants will hone skills in data preprocessing, model evaluation, and understanding trade-offs, essential in the data science realm.
- The case emphasizes actionable insights, fostering the ability to drive data-informed strategies in financial sectors.

Dataset Explanation: LoanTapData.csv (Link:

https://drive.google.com/file/d/1ZPYj7CZCfxntE8p2Lze_4QO4MyEOy6_d/view?usp=sharing)

1. loan_amnt: Amount borrower applied for.
2. term: Loan duration (36 or 60 months).
3. int_rate: Interest rate on loan.
4. installment: Monthly repayment amount.
5. grade: LoanTap assigned loan grade (Risk ratings by LoanTap.)
6. sub_grade: LoanTap assigned loan grade (Risk ratings by LoanTap.)
7. emp_title: Borrower's job title.
8. emp_length: Duration of borrower's employment (0-10 years).
9. home_ownership: Borrower's housing situation (own, rent, etc.).
10. annual_inc: Borrower's yearly income.
11. verification_status: Whether borrower's income was verified.

12. issue_d: Loan issuance month.
13. loan_status: Current status of the loan.
14. purpose: Borrower's reason for the loan.
15. title: The loan's title provided by the borrower.
16. dti (Debt-to-Income ratio): Monthly debt vs. monthly income ratio.
17. earliest_cr_line: Date of borrower's oldest credit account.
18. open_acc: Number of borrower's active credit lines.
19. pub_rec: Negative records on borrower's public credit profile.
20. revol_bal: Total credit balance.
21. revol_util: Usage percentage of 'revolving' accounts like credit cards.
22. total_acc: Total number of borrower's credit lines.
23. initial_list_status: Loan's first category ('W' or 'F').
24. application_type: Individual or joint application.
25. mort_acc: Number of borrower's mortgages.
26. pub_rec_bankruptcies: Bankruptcy records for borrower.
27. Address: Borrower's location.

What is Expected?

Assuming you are a data scientist at LoanTap, you are tasked with analyzing the dataset to determine the creditworthiness of potential borrowers. Your ultimate objective is to build a logistic regression model, evaluate its performance, and provide actionable insights for the underwriting process.

##Initial Analysis

Importing libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.impute import SimpleImputer

from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import StandardScaler

!pip install category_encoders
import category_encoders as ce

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression

from sklearn.metrics import (
    confusion_matrix
    ,ConfusionMatrixDisplay
    ,accuracy_score
```

```

,precision_score
,recall_score
,f1_score
,roc_curve
,roc_auc_score
,precision_recall_curve
,auc
)

from statsmodels.stats.outliers_influence import
variance_inflation_factor

from sklearn.model_selection import KFold, cross_val_score

from imblearn.over_sampling import SMOTE

import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
pd.set_option('display.max_columns', None)

Requirement already satisfied: category_encoders in
/usr/local/lib/python3.10/dist-packages (2.6.3)
Requirement already satisfied: numpy>=1.14.0 in
/usr/local/lib/python3.10/dist-packages (from category_encoders)
(1.25.2)
Requirement already satisfied: scikit-learn>=0.20.0 in
/usr/local/lib/python3.10/dist-packages (from category_encoders)
(1.2.2)
Requirement already satisfied: scipy>=1.0.0 in
/usr/local/lib/python3.10/dist-packages (from category_encoders)
(1.11.4)
Requirement already satisfied: statsmodels>=0.9.0 in
/usr/local/lib/python3.10/dist-packages (from category_encoders)
(0.14.2)
Requirement already satisfied: pandas>=1.0.5 in
/usr/local/lib/python3.10/dist-packages (from category_encoders)
(2.0.3)
Requirement already satisfied: patsy>=0.5.1 in
/usr/local/lib/python3.10/dist-packages (from category_encoders)
(0.5.6)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.0.5-
>category_encoders) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.0.5-
>category_encoders) (2023.4)
Requirement already satisfied: tzdata>=2022.1 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.0.5-
>category_encoders) (2024.1)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-

```

```

packages (from patsy>=0.5.1->category_encoders) (1.16.0)
Requirement already satisfied: joblib>=1.1.1 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.20.0-
>category_encoders) (1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.20.0-
>category_encoders) (3.5.0)
Requirement already satisfied: packaging>=21.3 in
/usr/local/lib/python3.10/dist-packages (from statsmodels>=0.9.0-
>category_encoders) (24.0)

```

```
data = pd.read_csv('logistic_regression.csv')
```

```
data.shape
```

```
(396030, 27)
```

There are 396030 rows and 27 columns. Total no of features = 26 as 'Loan Status' is target variable which our model should be predicting. We have data about 396030 loan applications with their loan status which tells us whether borrower defaulted or not. Thus, it's supervised ML problem.

```
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):

```

#	Column	Non-Null Count	Dtype
0	loan_amnt	396030 non-null	float64
1	term	396030 non-null	object
2	int_rate	396030 non-null	float64
3	installment	396030 non-null	float64
4	grade	396030 non-null	object
5	sub_grade	396030 non-null	object
6	emp_title	373103 non-null	object
7	emp_length	377729 non-null	object
8	home_ownership	396030 non-null	object
9	annual_inc	396030 non-null	float64
10	verification_status	396030 non-null	object
11	issue_d	396030 non-null	object
12	loan_status	396030 non-null	object
13	purpose	396030 non-null	object
14	title	394274 non-null	object
15	dti	396030 non-null	float64
16	earliest_cr_line	396030 non-null	object
17	open_acc	396030 non-null	float64
18	pub_rec	396030 non-null	float64
19	revol_bal	396030 non-null	float64
20	revol_util	395754 non-null	float64

```
21 total_acc          396030 non-null float64
22 initial_list_status 396030 non-null object
23 application_type     396030 non-null object
24 mort_acc            358235 non-null float64
25 pub_rec_bankruptcies 395495 non-null float64
26 address             396030 non-null object
```

```
dtypes: float64(12), object(15)
```

```
memory usage: 81.6+ MB
```

```
data.head()
```

```
{"type": "dataframe", "variable_name": "data"}
```

```
data.columns
```

```
Index(['loan_amnt', 'term', 'int_rate', 'installment', 'grade',
      'sub_grade',
      'emp_title', 'emp_length', 'home_ownership', 'annual_inc',
      'verification_status', 'issue_d', 'loan_status', 'purpose',
      'title',
      'dti', 'earliest_cr_line', 'open_acc', 'pub_rec', 'revol_bal',
      'revol_util', 'total_acc', 'initial_list_status',
      'application_type',
      'mort_acc', 'pub_rec_bankruptcies', 'address'],
      dtype='object')
```

```
data.nunique()
```

```
loan_amnt          1397
term                2
int_rate           566
installment        55706
grade              7
sub_grade          35
emp_title          173105
emp_length         11
home_ownership      6
annual_inc         27197
verification_status 3
issue_d            115
loan_status         2
purpose            14
title              48816
dti                4262
earliest_cr_line   684
open_acc           61
pub_rec            20
revol_bal          55622
revol_util         1226
total_acc          118
initial_list_status 2
```

```
application_type      3
mort_acc              33
pub_rec_bankruptcies  9
address              393700
dtype: int64
```

```
cat_columns = ['term', 'grade', 'sub_grade', 'emp_length',
               'home_ownership', 'verification_status', 'loan_status', 'purpose',
               'initial_list_status', 'application_type', 'emp_title', 'title']
```

```
for col in cat_columns:
    print(data[col].value_counts())
```

```
term
36 months    302005
60 months     94025
Name: count, dtype: int64
```

```
grade
B    116018
C    105987
A     64187
D     63524
E     31488
F     11772
G      3054
Name: count, dtype: int64
```

```
sub_grade
B3    26655
B4    25601
C1    23662
C2    22580
B2    22495
B5    22085
C3    21221
C4    20280
B1    19182
A5    18526
C5    18244
D1    15993
A4    15789
D2    13951
D3    12223
D4    11657
A3    10576
A1     9729
D5     9700
A2     9567
E1     7917
E2     7431
E3     6207
```

E4	5361
E5	4572
F1	3536
F2	2766
F3	2286
F4	1787
F5	1397
G1	1058
G2	754
G3	552
G4	374
G5	316

Name: count, dtype: int64

emp_length	
10+ years	126041
2 years	35827
< 1 year	31725
3 years	31665
5 years	26495
1 year	25882
4 years	23952
6 years	20841
7 years	20819
8 years	19168
9 years	15314

Name: count, dtype: int64

home_ownership	
MORTGAGE	198348
RENT	159790
OWN	37746
OTHER	112
NONE	31
ANY	3

Name: count, dtype: int64

verification_status	
Verified	139563
Source Verified	131385
Not Verified	125082

Name: count, dtype: int64

loan_status	
Fully Paid	318357
Charged Off	77673

Name: count, dtype: int64

purpose	
debt_consolidation	234507
credit_card	83019
home_improvement	24030
other	21185
major_purchase	8790

```

small_business      5701
car                 4697
medical             4196
moving              2854
vacation            2452
house               2201
wedding             1812
renewable_energy    329
educational         257
Name: count, dtype: int64
initial_list_status
f      238066
w      157964
Name: count, dtype: int64
application_type
INDIVIDUAL      395319
JOINT            425
DIRECT_PAY       286
Name: count, dtype: int64
emp_title
Teacher          4389
Manager          4250
Registered Nurse  1856
RN               1846
Supervisor       1830
...
Postman          1
McCarthy & Holthus, LLC  1
jp flooring      1
Histology Technologist  1
Gracon Services, Inc  1
Name: count, Length: 173105, dtype: int64
title
Debt consolidation      152472
Credit card refinancing  51487
Home improvement        15264
Other                   12930
Debt Consolidation      11608
...
Graduation/Travel Expenses  1
Daughter's Wedding Bill    1
gotta move                 1
creditcardrefi             1
Toxic Debt Payoff          1
Name: count, Length: 48816, dtype: int64

data.describe(include='all')

{"type": "dataframe"}

```


1. Loan Amount has a range of [500, 40000] with median being 12000.
2. The term for the loan is usually 36 months.
3. Median interest rate on the loan is 13.3%.
4. Monthly repayment amount is having median value of ~ 375.
5. Most of the borrowers fall under loan grade 'B' and sub-grade 'B3' (Risk rating by LoanTap)
6. The most common borrower's job title is 'Teacher' and duration of employment is '10+ years'.
7. Borrower's housing situation is usually 'Mortgage'.
8. Median Annual Income is 64000. And, it's verified in the most cases as it's necessary before approving the loan.
9. Loan issuance month is Oct-2014 for most of the data-points here.
10. Loan status is 'Fully Paid' in most cases here which means non-defaulter is majority which is great for the company.
11. In most loan application cases, borrower's reason for the loan and the loan's title provided by the borrower is 'Debt consolidation'.
12. The dti (Monthly debt vs. monthly income ratio) has median value of 16.91
13. Date of borrower's oldest credit account is Oct-2000 in most cases.
14. Number of borrower's active credit lines has median value of 10.
15. Negative records on borrower's public credit profile has median of 0 which is superb.
16. Total credit balance has median value of ~11,000
17. Usage percentage of 'revolving' accounts like credit cards has median of 54.8
18. Total number of borrower's credit lines has median of 24.
19. Loan's first category has value 'f' in most cases.
20. Application type is mostly 'Individual'.
21. Number of borrower's mortgages has median value 1.
22. Bankruptcy records for borrower is 0 in 50% of cases.
23. Borrower's location is 'AE 70466' in the most cases.

Data Exploration

```
data.groupby(by = 'loan_status')['loan_amnt'].describe()

{"summary": "{\n  \"name\": \"data\",\n  \"rows\": 2,\n  \"fields\": [\n    {\n      \"column\": \"loan_status\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 2,\n        \"samples\": [\n          \"Fully Paid\",\n          \"Charged Off\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"count\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 170189.288523103,\n        \"min\": 77673.0,\n        \"max\": 318357.0,\n        \"num_unique_values\": 2,\n        \"samples\": [\n          318357.0,\n          77673.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ],\n  \"column\": \"mean\",\n  \"properties\": {\n    \"dtype\": \"number\",\n    \"std\": 890.5459748555348,\n    \"min\":
```

```

13866.878771316478,\n          \"max\": 15126.300966873945,\n
\"num_unique_values\": 2,\n          \"samples\": [\n
13866.878771316478,\n          15126.300966873945\n          ],\n
\"semantic_type\": \"\", \n          \"description\": \"\"\n
n    },\n    {\n          \"column\": \"std\", \n          \"properties\": {\n
\"dtype\": \"number\", \n          \"std\": 143.38064827557537,\n
\"min\": 8302.319699344323,\n          \"max\": 8505.090556717489,\n
\"num_unique_values\": 2,\n          \"samples\": [\n
8302.319699344323,\n          8505.090556717489\n          ],\n
\"semantic_type\": \"\", \n          \"description\": \"\"\n
n    },\n    {\n          \"column\": \"min\", \n          \"properties\": {\n
\"dtype\": \"number\", \n          \"std\": 353.5533905932738,\n
\"min\": 500.0,\n          \"max\": 1000.0,\n
\"num_unique_values\": 2,\n          \"samples\": [\n          500.0,\n
1000.0\n          ],\n          \"semantic_type\": \"\", \n          \"description\": \"\"\n
n    },\n    {\n          \"column\": \"25%\", \n          \"properties\": {\n          \"dtype\": \"number\", \n
\"std\": 724.7844507162113,\n          \"min\": 7500.0,\n
\"max\": 8525.0,\n          \"num_unique_values\": 2,\n
\"samples\": [\n          7500.0,\n          8525.0\n          ],\n
\"semantic_type\": \"\", \n          \"description\": \"\"\n
n    },\n    {\n          \"column\": \"50%\", \n          \"properties\": {\n
\"dtype\": \"number\", \n          \"std\": 1414.213562373095,\n
\"min\": 12000.0,\n          \"max\": 14000.0,\n
\"num_unique_values\": 2,\n          \"samples\": [\n          12000.0,\n
14000.0\n          ],\n          \"semantic_type\": \"\", \n          \"description\": \"\"\n
n    },\n    {\n          \"column\": \"75%\", \n          \"properties\": {\n          \"dtype\": \"number\", \n
\"std\": 548.0077554195743,\n          \"min\": 19225.0,\n
\"max\": 20000.0,\n          \"num_unique_values\": 2,\n
\"samples\": [\n          19225.0,\n          20000.0\n          ],\n
\"semantic_type\": \"\", \n          \"description\": \"\"\n
n    },\n    {\n          \"column\": \"max\", \n          \"properties\": {\n
\"dtype\": \"number\", \n          \"std\": 0.0,\n          \"min\": 40000.0,\n
\"max\": 40000.0,\n          \"num_unique_values\": 1,\n
\"samples\": [\n          40000.0\n          ],\n
\"semantic_type\": \"\", \n          \"description\": \"\"\n
n    }\n  ]\n}, \"type\": \"dataframe\"}

```

The no of people those who have fully paid are 318357 and that of Charged Off are 77673.

```
data['home_ownership'].value_counts()
```

```

home_ownership
MORTGAGE      198348
RENT          159790
OWN           37746
OTHER         112
NONE          31

```

```
ANY          3
Name: count, dtype: int64
```

The majority of people have home ownership as Mortgage and Rent.

Combining the minority classes as 'OTHER'.

```
data.loc[(data['home_ownership'] == 'ANY') | (data['home_ownership']
== 'NONE'), 'home_ownership'] = 'OTHER'
data['home_ownership'].value_counts()
```

```
home_ownership
MORTGAGE      198348
RENT          159790
OWN           37746
OTHER         146
Name: count, dtype: int64
```

#Checking the distribution of OTHER

```
data.loc[data['home_ownership'] == 'OTHER',
'loan_status'].value_counts()
```

```
loan_status
Fully Paid      123
Charged Off     23
Name: count, dtype: int64
```

Issues in title, looks like values were manually entered

```
data['title'].value_counts()[:20]
```

```
title
Debt consolidation      152472
Credit card refinancing  51487
Home improvement        15264
Other                   12930
Debt Consolidation      11608
Major purchase          4769
Consolidation           3852
debt consolidation      3547
Business                2949
Debt Consolidation Loan  2864
Medical expenses        2742
Car financing           2139
Credit Card Consolidation 1775
Vacation                1717
Moving and relocation    1689
consolidation           1595
Personal Loan           1591
```

```

Consolidation Loan      1299
Home Improvement        1268
Home buying             1183
Name: count, dtype: int64

data['title'] = data.title.str.lower()

data['title'].value_counts()[:10]

title
debt consolidation      168108
credit card refinancing  51781
home improvement        17117
other                   12993
consolidation           5583
major purchase          4998
debt consolidation loan  3513
business                3017
medical expenses        2820
credit card consolidation 2638
Name: count, dtype: int64

```



##Univariate Analysis



###Continuous Variables



```



for col in data.columns:
    if data[col].dtype in ('float64', 'int64'):
        sns.boxplot(data = data, x = col, palette = 'ocean')
        plt.show()



```







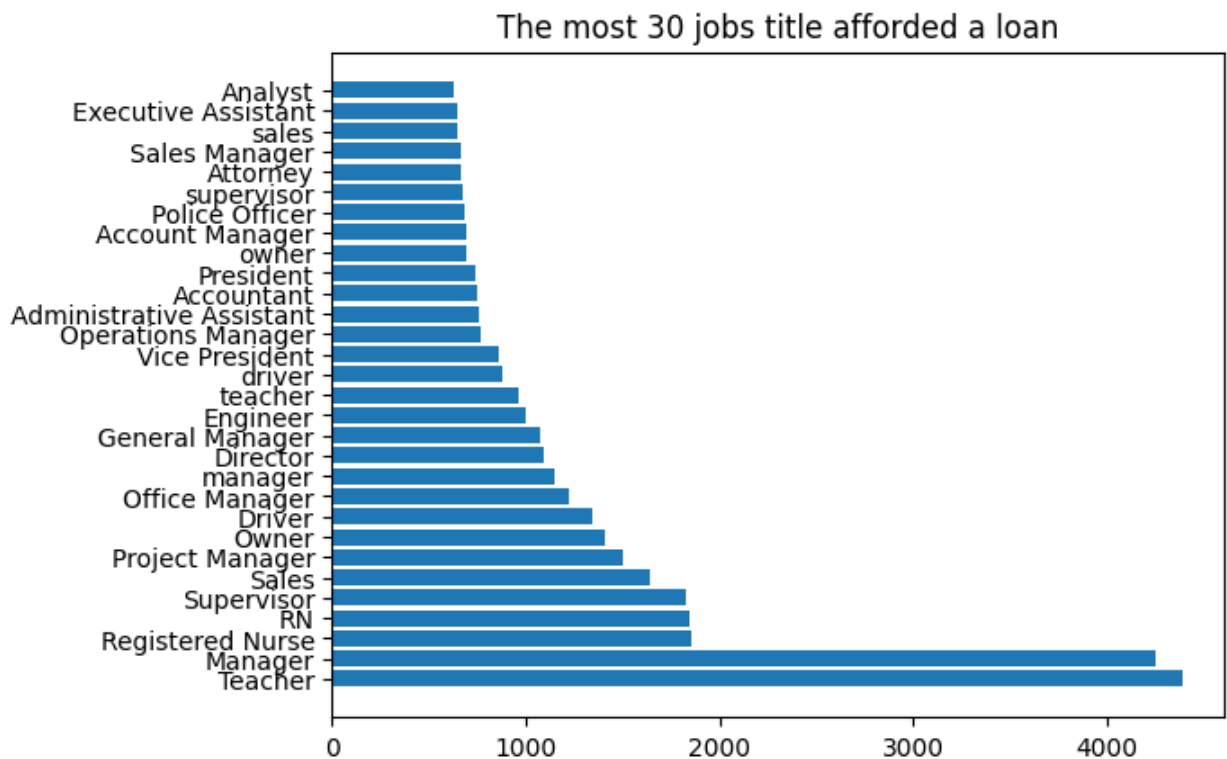




The above box-plots tell us that there are outliers present in all the continuous variables.

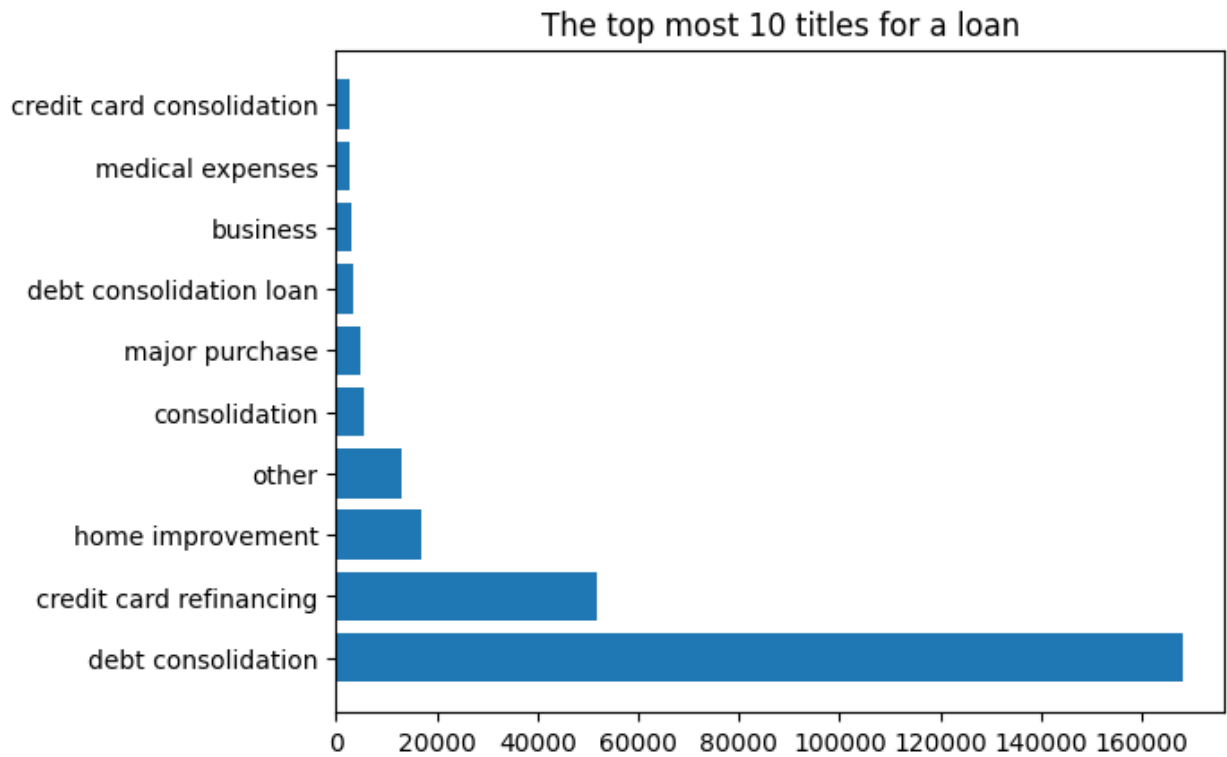
###Categorical Variables

```
plt.barh(data.emp_title.value_counts()[:30].index,  
data.emp_title.value_counts()[:30])  
plt.title("The most 30 jobs title afforded a loan")  
plt.show()
```





The top 2 employee job roles for which we have the highest no of loan applications are 'Teacher' followed by 'Manager'.



```
plt.barh(data.title.value_counts()[:10].index,  
data.title.value_counts()[:10])  
plt.title("The top most 10 titles for a loan")  
plt.show()
```







The loan's title provided by the borrower is 'Debt Consolidation' for the most no of applications.



```
for col in cat_columns:  
    if col not in ('emp_title', 'title'):  
        plt.barh(data[col].value_counts().index, data[col].value_counts())  
        plt.ylabel(col)  
        plt.show()
```











1. We have majority of loan applications for the loan term of '36 months'. Which tells us that borrowers prefer lesser duration for the loan term.

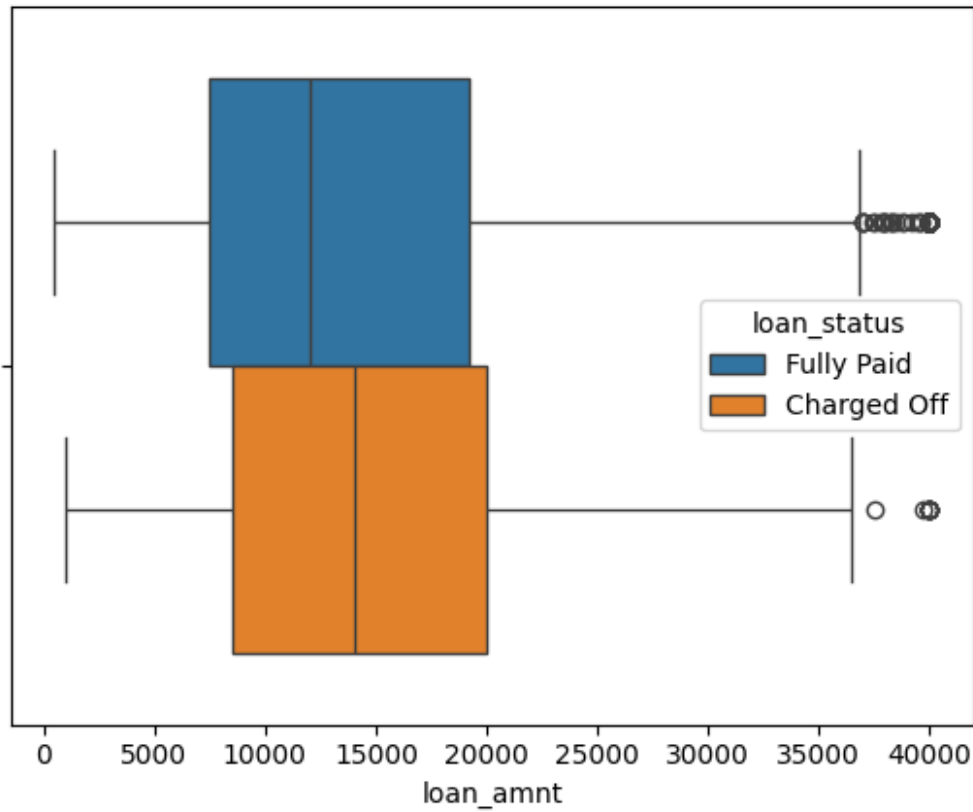
2. The loan applications fall under the grades 'B' followed by 'C', and sub-grade 'B3' followed by 'B4'.
3. Majority of the borrowers have the employment duration more than 10 years.
4. Top home ownership options observed are 'Mortgage' followed by 'Rent'.
5. The no of loan applications among different borrower income verification status ('Not Verified', 'Source Verified' and 'Verified') is almost similar which is not good as it's important to verify the income before approving the loan to minimize the defaulters and the subsequent money loss.
6. Majority of the loan applications are fully paid which is good.
7. Top 2 purposes for which loan was taken are 'Debt Consolidation' and 'Credit Card'.
8. The initial list status for majority of loan applications is 'f' which means fractional.
9. Loan application type is mostly 'Individual'.

##Bivariate Analysis


Target (Categorical) vs Features (Continuous)

```
for col in data.columns:  
    if data[col].dtype != 'object':  
        sns.boxplot(data = data, x = col, hue = 'loan_status', orient =  
'v')  
        plt.show()
```

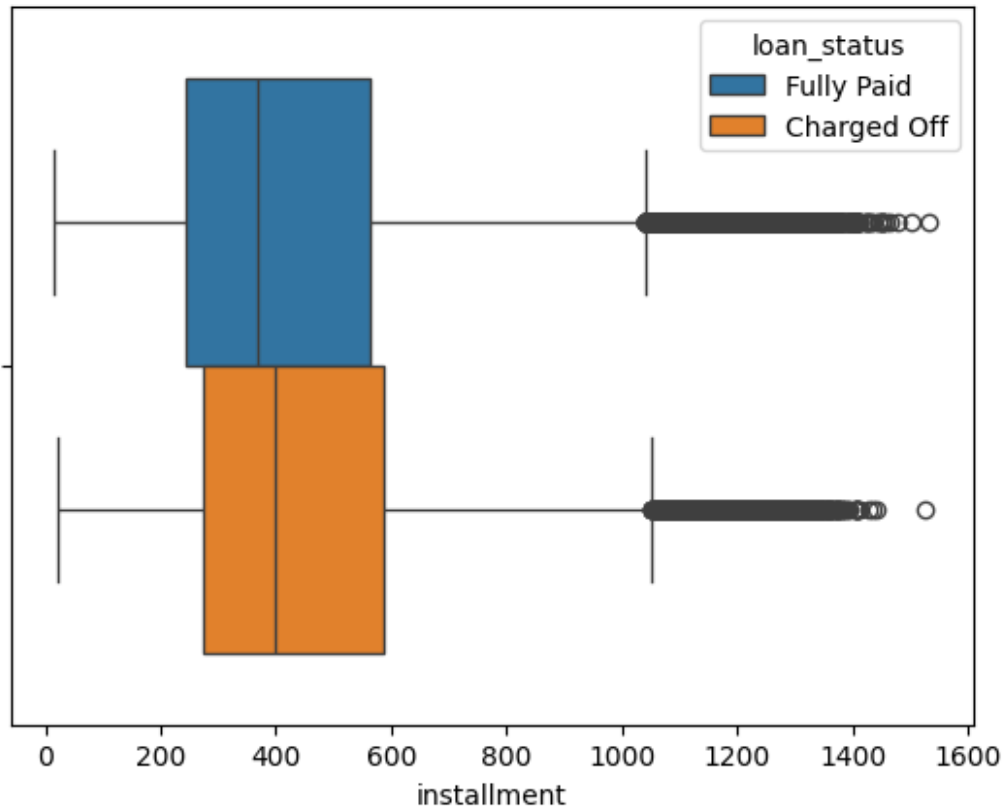
```
/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:1608:  
UserWarning: Vertical orientation ignored with only `x` specified.  
warnings.warn(single_var_warning.format("Vertical", "x"))
```




```
/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:1608:  
UserWarning: Vertical orientation ignored with only `x` specified.  
warnings.warn(single_var_warning.format("Vertical", "x"))
```



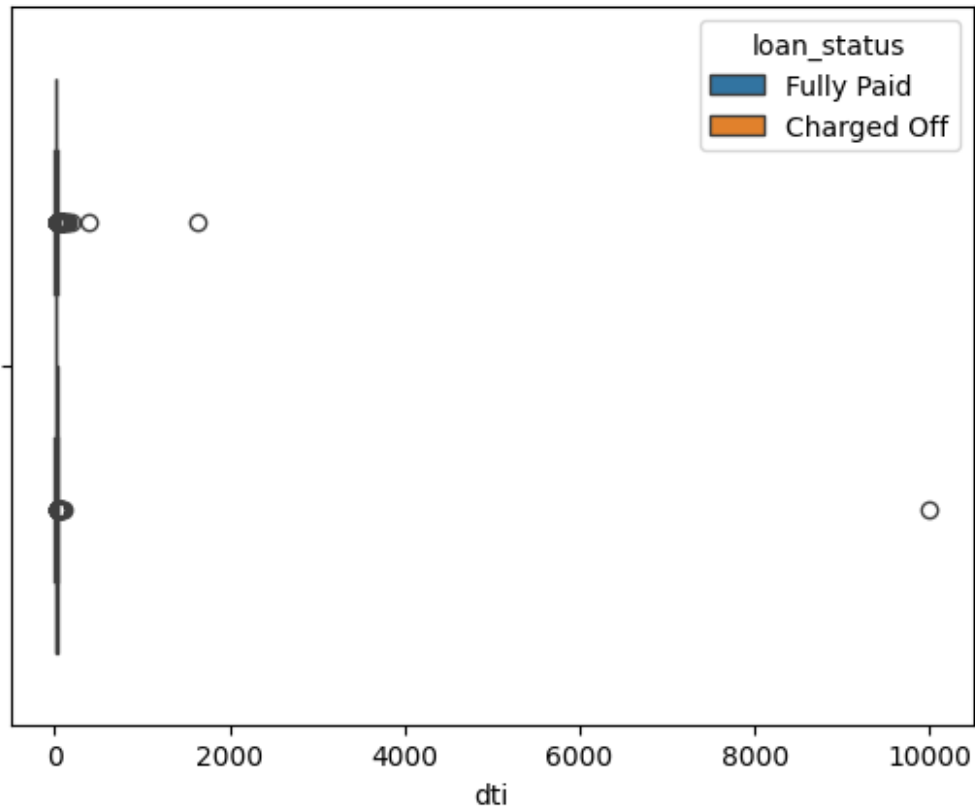
```
/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:1608:  
UserWarning: Vertical orientation ignored with only `x` specified.  
warnings.warn(single_var_warning.format("Vertical", "x"))
```




```
/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:1608:  
UserWarning: Vertical orientation ignored with only `x` specified.  
warnings.warn(single_var_warning.format("Vertical", "x"))
```




```
/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:1608:
UserWarning: Vertical orientation ignored with only `x` specified.
warnings.warn(single_var_warning.format("Vertical", "x"))
```



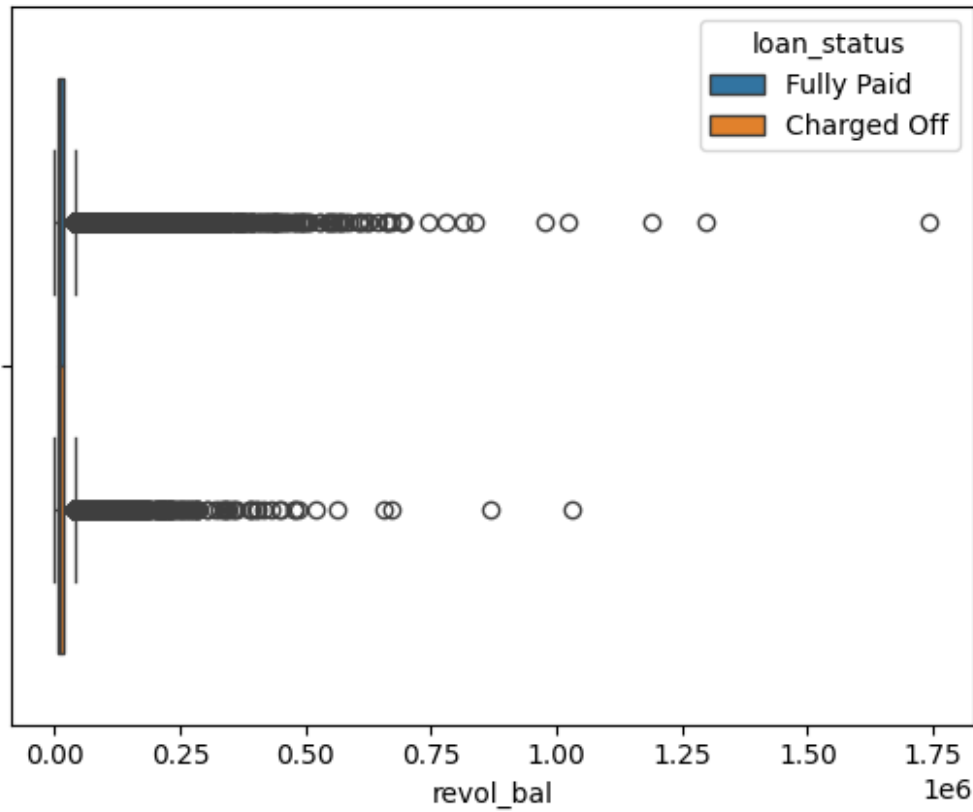
```
/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:1608:  
UserWarning: Vertical orientation ignored with only `x` specified.  
warnings.warn(single_var_warning.format("Vertical", "x"))
```




```
/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:1608:  
UserWarning: Vertical orientation ignored with only `x` specified.  
warnings.warn(single_var_warning.format("Vertical", "x"))
```


```
/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:1608:  
UserWarning: Vertical orientation ignored with only `x` specified.  
warnings.warn(single_var_warning.format("Vertical", "x"))
```




```
/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:1608:  
UserWarning: Vertical orientation ignored with only `x` specified.  
warnings.warn(single_var_warning.format("Vertical", "x"))
```




```
/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:1608:  
UserWarning: Vertical orientation ignored with only `x` specified.  
warnings.warn(single_var_warning.format("Vertical", "x"))
```



```
/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:1608:  
UserWarning: Vertical orientation ignored with only `x` specified.  
warnings.warn(single_var_warning.format("Vertical", "x"))
```




```
/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:1608:  
UserWarning: Vertical orientation ignored with only `x` specified.  
warnings.warn(single_var_warning.format("Vertical", "x"))
```






The median values for loan amount, installment, interest rate, no of active credit lines are slightly greater for the defaulters which makes sense as they couldn't pay off the loan.


Target (Categorical) vs Features (Categorical)


```
for col in cat_columns:
    if col not in ('emp_title', 'title'):
        sns.countplot(data = data, x = col, fill = 'True', hue =
'loan_status')
        plt.xticks(rotation = 90)
        plt.show()
```
























1. For loan term '60 months', there is no huge difference between the no of defaulters and non-defaulters which means for greater duration the risk of defaulting is also high.
2. For the loan ratings 'A', 'B', 'C' and 'D', there's a huge difference between the no of defaulters and non-defaulters and no of non-defaulters is high which means more likely loan will be paid but for grades 'E', 'F' and 'G', these nos are almost same that means high risk ratings.
3. Thus, we can say that grades 'A', 'B' and 'C' are low risk, grades 'D' and 'E' are moderate risk and 'F' and 'G' are high risk.
4. Similar pattern observed for the sub-grades with 1 being low risk in that grade and 5 being high risk.
5. So overall, A1 is lowest risk and G5 is the highest risk.
6. Among the borrowers with initial list status as 'w' i.e. whole is difference is lesser as compared to 'f' which means that more chance of defaulting when entire amount is approved.
7. Employee tenure, income verification status, purpose of loan, home ownership status and loan application type do not make much impact on defaulting.

Features : Continuous vs Continuous

```
num_cols = [col for col in data.columns if data[col].dtype !=  
'object']  
  
fig, ax = plt.subplots(figsize=(10,10))  
sns.heatmap(data[num_cols].corr(method = 'spearman'), annot = True, ax  
= ax)  
plt.show()
```



1. The correlation coefficient is the highest for loan amount and installment features which is +0.97 which tells us that greater the loan amount, greater will be the monthly installment amount.
2. Also, the correlation is quite high (+0.86) between pub_rec (Negative records on borrower's public credit profile) and pub_rec_bankruptcies (Bankruptcy records for borrower) which is justified as the no of bankruptcies increases, the no of negative records would also increase.
3. Another pair having higher correlation (+0.67) is open_acc (Number of borrower's active credit lines) and total_acc (Total number of borrower's credit lines)
4. There's negative correlation (-0.1) between loan amount and pub_rec_bankruptcies or pub_rec as no of negative records or bankruptcies increases loan amount would decrease and it becomes quite risky to assume that borrower won't default.
5. Also, annual income and dti have negative correlation (-0.2) as monthly debt to monthly income ratio would decrease with increase in the income.

We noticed almost perfect correlation between "loan_amnt" the "installment" feature.

- installment: The monthly payment owed by the borrower if the loan originates.
- loan_amnt: The listed amount of the loan applied for by the borrower. If at some point in time, the credit department reduces the loan amount, then it will be reflected in this value.

So, we can drop either one of those columns.

```
data.drop(columns=['installment'], axis = 1, inplace = True)
```

#Data Preprocessing

Duplicate value check

```
data.duplicated().sum()
```

```
0
```

There are no duplicate rows present in the dataset.

Missing Value Detection

```
data.isna().sum()
```

loan_amnt	0
term	0
int_rate	0
grade	0
sub_grade	0
emp_title	22927
emp_length	18301
home_ownership	0
annual_inc	0

```

verification_status      0
issue_d                  0
loan_status              0
purpose                 0
title                   1756
dti                     0
earliest_cr_line         0
open_acc                0
pub_rec                 0
revol_bal               0
revol_util              276
total_acc               0
initial_list_status      0
application_type         0
mort_acc               37795
pub_rec_bankruptcies    535
address                 0
dtype: int64

data.isna().sum() * 100.0/len(data)

loan_amnt      0.000000
term           0.000000
int_rate       0.000000
grade          0.000000
sub_grade      0.000000
emp_title      5.789208
emp_length     4.621115
home_ownership 0.000000
annual_inc     0.000000
verification_status 0.000000
issue_d        0.000000
loan_status    0.000000
purpose        0.000000
title          0.443401
dti            0.000000
earliest_cr_line 0.000000
open_acc       0.000000
pub_rec        0.000000
revol_bal      0.000000
revol_util     0.069692
total_acc      0.000000
initial_list_status 0.000000
application_type 0.000000
mort_acc       9.543469
pub_rec_bankruptcies 0.135091
address        0.000000
dtype: float64

```

There are missing values present in the columns:

1. emp_title
2. emp_length
3. title
4. revol_util
5. mort_acc
6. pub_rec_bankruptcies

Missing Value Treatment

We will be using Simple Imputer to fill the missing values by using 'Most Frequent' strategy for the categorical variables such as emp_title, title and emp_length.

```
imp = SimpleImputer(missing_values=np.nan, strategy='most_frequent')
data[['emp_title_filled', 'emp_length_filled', 'title_filled']] =
imp.fit_transform(data[['emp_title', 'emp_length', 'title']])
```

Checking if missing values got filled correctly for categorical variables `emp_title`, `title` and `emp_length`.

```
data[['emp_title']].mode()

{"summary": "{\n  \"name\": \"data[['emp_title']]\", \n  \"rows\": 1,\n  \"fields\": [\n    {\n      \"column\": \"emp_title\", \n      \"dtype\": \"string\", \n      \"num_unique_values\": 1,\n      \"samples\": [\n        \"Teacher\", \n        \"\", \n        \"\" \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  ] \n}", "type": "dataframe"}

data[['emp_length']].mode()

{"summary": "{\n  \"name\": \"data[['emp_length']]\", \n  \"rows\": 1,\n  \"fields\": [\n    {\n      \"column\": \"emp_length\", \n      \"dtype\": \"string\", \n      \"num_unique_values\": 1,\n      \"samples\": [\n        \"10+ years\", \n        \"\" \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  ] \n}", "type": "dataframe"}

data[['title']].mode()

{"summary": "{\n  \"name\": \"data[['title']]\", \n  \"rows\": 1,\n  \"fields\": [\n    {\n      \"column\": \"title\", \n      \"dtype\": \"string\", \n      \"num_unique_values\": 1,\n      \"samples\": [\n        \"debt consolidation\", \n        \"\" \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  ] \n}", "type": "dataframe"}

data[data[['title']].isna().any(axis = 1)][['emp_title', 'emp_title_filled', 'emp_length', 'emp_length_filled', 'title', 'title_filled']]
```

```
{"repr_error": "0", "type": "dataframe"}
```

Using strategy 'Median' for the continuous variables like `revol_util`, `mort_acc` and `pub_rec_bankruptcies`.

```
imp = SimpleImputer(missing_values=np.nan, strategy='median')
data[['mort_acc_filled', 'revol_util_filled',
      'pub_rec_bankruptcies_filled']] = imp.fit_transform(data[['mort_acc',
      'revol_util', 'pub_rec_bankruptcies']])

data[['mort_acc']].median()

mort_acc    1.0
dtype: float64

data[['revol_util']].median()

revol_util   54.8
dtype: float64

data[['pub_rec_bankruptcies']].median()

pub_rec_bankruptcies    0.0
dtype: float64
```

Drop original columns for these 6 that we filled and rename the `_filled` ones with the original names.

```
data.drop(['emp_title', 'emp_length', 'title', 'mort_acc',
          'revol_util', 'pub_rec_bankruptcies'], axis = 1, inplace = True)

data.isna().sum()

loan_amnt    0
term         0
int_rate     0
grade        0
sub_grade    0
home_ownership  0
annual_inc   0
verification_status  0
issue_d      0
loan_status  0
purpose      0
dti          0
earliest_cr_line  0
open_acc     0
pub_rec      0
revol_bal    0
total_acc    0
```

```

initial_list_status      0
application_type         0
address                 0
emp_title_filled         0
emp_length_filled        0
title_filled            0
mort_acc_filled         0
revol_util_filled        0
pub_rec_bankruptcies_filled 0
dtype: int64

data = data.rename(columns = {'emp_title_filled': 'emp_title',
                              'emp_length_filled' : 'emp_length', 'title_filled' : 'title',
                              'mort_acc_filled' : 'mort_acc', 'revol_util_filled' : 'revol_util',
                              'pub_rec_bankruptcies_filled' : 'pub_rec_bankruptcies'})

```

Outlier Detection and Treatment

Using IQR method to detect and treat the outliers.

```

#Calculating few more statistical measures such as 'Range', 'IQR',
'Lower Whisker' and 'Upper Whisker'

descriptive_stats = data.describe()
descriptive_stats =
descriptive_stats.reindex(descriptive_stats.index.values.tolist()+
['Range', 'IQR', 'Lower Whisker', 'Upper Whisker'])

for col in descriptive_stats.columns:
    descriptive_stats.loc['Range'][col] = descriptive_stats.loc['max']
[col] - descriptive_stats.loc['min'][col]
    descriptive_stats.loc['IQR'][col] = descriptive_stats.loc['75%']
[col] - descriptive_stats.loc['25%'][col]
    descriptive_stats.loc['Lower Whisker'][col] =
descriptive_stats.loc['25%'][col] - (1.5 *
descriptive_stats.loc['IQR'][col])
    descriptive_stats.loc['Upper Whisker'][col] =
descriptive_stats.loc['75%'][col] + (1.5 *
descriptive_stats.loc['IQR'][col])

descriptive_stats

{"summary": "{\n  \"name\": \"descriptive_stats\", \n  \"rows\": 12, \n
  \"fields\": [\n    {\n      \"column\": \"loan_amnt\", \n
  \"properties\": {\n      \"dtype\": \"number\", \n      \"std\":
110649.59456444613, \n      \"min\": -10000.0, \n      \"max\":
396030.0, \n      \"num_unique_values\": 11, \n      \"samples\": [\n
12000.0, \n      396030.0, \n      -10000.0\
n      ], \n      \"semantic_type\": \"\", \n

```

```

{"description": "", "int_rate": 1.49, "number": 1.49, "num_unique_values": 12, "properties": {"dtype": "number", "std": 114319.98871432777, "min": 1.49, "max": 396030.0, "samples": [1.49, 5.9999999999999998, 396030.0]}, "semantic_type": "number", "column": "int_rate"}, {"description": "annual_inc", "number": 74203.17579771738, "num_unique_values": 10, "properties": {"dtype": "number", "std": 3355267.087918833, "min": -22500.0, "max": 8706582.0, "samples": [-22500.0, 74203.17579771738, 64000.0]}, "semantic_type": "number", "column": "annual_inc"}, {"description": "open_acc", "number": 114317.28801162555, "num_unique_values": 11, "properties": {"dtype": "number", "std": 113860.8635870796, "min": -6.27, "max": 396030.0, "samples": [-6.27, 16.91, 396030.0]}, "semantic_type": "number", "column": "open_acc"}, {"description": "pub_rec", "number": 114319.48592479543, "num_unique_values": 5, "properties": {"dtype": "number", "std": 114319.48592479543, "min": 0.0, "max": 396030.0, "samples": [0.0, 86.0, 0.530670600474012]}, "semantic_type": "number", "column": "pub_rec"}, {"description": "revol_bal", "number": 667965.7781961308, "num_unique_values": 11, "properties": {"dtype": "number", "std": 1743266.0, "min": -14367.5, "max": 396030.0, "samples": [-14367.5, 11181.0, 396030.0]}, "semantic_type": "number", "column": "revol_bal"}, {"description": "total_acc", "number": 114311.52591669752, "num_unique_values": 12, "properties": {"dtype": "number", "std": 114311.52591669752, "min": -5.5, "max": 396030.0, "samples": [-5.5, 15.0, 396030.0]}, "semantic_type": "number", "column": "total_acc"}, {"description": "mort_acc", "number": 114321.86775896826, "num_unique_values": 9, "properties": {"dtype": "number", "std": 114321.86775896826, "min": -4.5, "max": 396030.0, "samples": [-4.5, 1.736307855465495, 3.0]}, "semantic_type": "number", "column": "mort_acc"}

```

```

{"semantic_type": "\n", "description": "\n", "column": "revol_util", "dtype": "number", "std": 114267.48240899428, "min": -19.600000000000016, "max": 396030.0, "num_unique_values": 11, "samples": [54.8, 396030.0, -19.600000000000016], "semantic_type": "\n", "description": "\n", "column": "pub_rec_bankruptcies", "dtype": "number", "std": 114323.58117393513, "min": 0.0, "max": 396030.0, "num_unique_values": 5, "samples": [0.12148322096810847, 8.0, 0.35596165879827396], "semantic_type": "\n", "description": "\n", "column": "dti", "dtype": "number", "std": 114267.48240899428, "min": -19.600000000000016, "max": 396030.0, "num_unique_values": 11, "samples": [54.8, 396030.0, -19.600000000000016]}
{"type": "dataframe", "variable_name": "descriptive_stats"}

```

No of outliers present in the respective column

```

for col in descriptive_stats.columns:
    if col not in cat_columns:
        print(col, ':', data[(data[col] < descriptive_stats.loc['Lower Whisker'][col]) | (data[col] > descriptive_stats.loc['Upper Whisker'][col])][col].count())

```

```

loan_amnt : 191
int_rate : 3777
annual_inc : 16700
dti : 275
open_acc : 10307
pub_rec : 57758
revol_bal : 21259
total_acc : 8499
mort_acc : 6843
revol_util : 12
pub_rec_bankruptcies : 45115

```

```

for col in descriptive_stats.columns:
    if col not in cat_columns:
        print(col, ':', data[(data[col] < descriptive_stats.loc['Lower Whisker'][col]) | (data[col] > descriptive_stats.loc['Upper Whisker'][col])][col].count()/len(data))

```

```

loan_amnt : 0.0004822866954523647
int_rate : 0.009537156276039694
annual_inc : 0.042168522586672724
dti : 0.0006943918390020958
open_acc : 0.02602580612579855
pub_rec : 0.145842486680302
revol_bal : 0.05368027674671111
total_acc : 0.021460495417013864

```

```
mort_acc : 0.017278994015604877
revol_util : 3.0300734792818725e-05
pub_rec_bankruptcies : 0.11391813751483473
```

We can see that pub_rec and pub_rec_bankruptcies have highest % of outliers among others.

```
for col in data.columns:
    if (col not in cat_columns) and (col not in
descriptive_stats.columns):
        print(col)
```

```
issue_d
earliest_cr_line
address
```

Clipping outliers to LW if < min and UW if > max.

```
for col in descriptive_stats.columns:
    if col not in cat_columns and col not in ('pub_rec', 'mort_acc',
'pub_rec_bankruptcies'):
        min = descriptive_stats.loc['Lower Whisker'][col]
        max = descriptive_stats.loc['Upper Whisker'][col]

        data.loc[data[col] < descriptive_stats.loc['Lower Whisker'][col],
col] = min
        data.loc[data[col] > descriptive_stats.loc['Upper Whisker'][col],
col] = max
```


Not dealing with the outliers present in pub_rec, mort_acc and pub_rec_bankruptcies as we are going to convert them to categorical columns by creating flags based on condition if > 1 then 1 else 0.



```
for col in descriptive_stats.columns:
    if col not in cat_columns:
        print(col, ': ', data[(data[col] < descriptive_stats.loc['Lower
Whisker'][col]) | (data[col] > descriptive_stats.loc['Upper Whisker']
[col])][col].count())
```



```
loan_amnt : 0
int_rate : 0
annual_inc : 0
dti : 0
open_acc : 0
pub_rec : 57758
revol_bal : 0
total_acc : 0
mort_acc : 6843
revol_util : 0
pub_rec_bankruptcies : 45115
```







```
for col in data.columns:  
    if data[col].dtype != 'object':  
        sns.boxplot(data = data, x = col, palette = 'ocean')  
        plt.show()
```















```
data[cat_columns]
{"type": "dataframe"}
```

Feature Engineering

New Feature: is_defaulter

```
data.loc[data['loan_status'] == 'Charged Off', 'is_defaulter'] = 1
data.loc[data['loan_status'] == 'Fully Paid', 'is_defaulter'] = 0

data.head()

{"type": "dataframe", "variable_name": "data"}

data['is_defaulter'].value_counts()

is_defaulter
0.0    318357
1.0     77673
Name: count, dtype: int64

data['is_defaulter'].value_counts(normalize = True)

is_defaulter
0.0    0.803871
1.0    0.196129
Name: proportion, dtype: float64
```

We can see that 80.4% are non-defaulters and 19.6% are defaulters.

Convert emp_length to numerical

```
data['emp_length'].value_counts()

emp_length
10+ years    144342
2 years      35827
< 1 year    31725
3 years      31665
5 years      26495
1 year       25882
4 years      23952
6 years      20841
7 years      20819
8 years      19168
9 years      15314
Name: count, dtype: int64
```

```

data['emp_length_num'] = (((data['emp_length'].str.replace('year',
'')).str.replace('s', '').str.strip()).replace('10+', 10)).replace('<
1', 0)).astype('float')

data[['emp_length', 'emp_length_num']]

{"type": "dataframe"}

data['emp_length_num'].value_counts()

emp_length_num
10.0      144342
2.0       35827
0.0       31725
3.0       31665
5.0       26495
1.0       25882
4.0       23952
6.0       20841
7.0       20819
8.0       19168
9.0       15314
Name: count, dtype: int64

data = data.drop('emp_length', axis = 1)
data = data.rename(columns = {'emp_length_num' : 'emp_length'})
data.head()

{"type": "dataframe", "variable_name": "data"}

```

Extract month and year from issue_d and earliest_cr_line

```

data[['issue_d_month', 'issue_d_year']] =
data['issue_d'].str.split('-', expand = True)
data[['earliest_cr_line_month', 'earliest_cr_line_year']] =
data['earliest_cr_line'].str.split('-', expand = True)

data.head()

{"type": "dataframe", "variable_name": "data"}

data['issue_d_month'].value_counts()

issue_d_month
Oct      42130
Jul       39714
Jan       34682
Nov       34068
Apr       33223
Aug       32816
Mar       31919

```



```
May      31895
Jun       30140
Dec       29082
Feb       28742
Sep       27619
Name: count, dtype: int64
```

```
data['issue_d_year'].value_counts()
```

```
issue_d_year
2014      102860
2013       97662
2015       94264
2012       41202
2016       28088
2011       17435
2010        9258
2009        3826
2008        1240
2007         195
Name: count, dtype: int64
```

```
data['earliest_cr_line_month'].value_counts()
```

```
earliest_cr_line_month
Oct       38291
Sep       37673
Aug       37349
Nov       35583
Dec       33687
Jul       31972
Mar       31617
Jan       30694
Jun       30445
May       30445
Apr       29231
Feb       29043
Name: count, dtype: int64
```

```
data['earliest_cr_line_year'].value_counts()
```

```
earliest_cr_line_year
2000       29366
2001       29083
1999       26491
2002       25901
2003       23657
...
1951         3
1950         3
1953         2
```

```
1944      1
1948      1
Name: count, Length: 65, dtype: int64
```

1. We can see that we have max loan applications from month of Oct and the year 2014.
2. First Credit line is mostly from year 2000.

```
data_dict = {'Jan':1, 'Feb':2, 'Mar':3, 'Apr':4, 'May':5, 'Jun':6,
             'Jul':7, 'Aug':8, 'Sep':9, 'Oct':10, 'Nov':11, 'Dec':12}
```

```
data['issue_d_month_no'] = data['issue_d_month'].map(data_dict)
```

```
data['earliest_cr_line_month_no'] =
data['earliest_cr_line_month'].map(data_dict)
```

```
data.head()
```

```
{"type":"dataframe","variable_name":"data"}
```

```
data['issue_d_month_no'].value_counts()
```

```
issue_d_month_no
```

```
10      42130
```

```
7       39714
```

```
1       34682
```

```
11      34068
```

```
4       33223
```

```
8       32816
```

```
3       31919
```

```
5       31895
```

```
6       30140
```

```
12      29082
```

```
2       28742
```

```
9       27619
```

```
Name: count, dtype: int64
```

```
data['earliest_cr_line_month_no'].value_counts()
```

```
earliest_cr_line_month_no
```

```
10      38291
```

```
9       37673
```

```
8       37349
```

```
11      35583
```

```
12      33687
```

```
7       31972
```

```
3       31617
```

```
1       30694
```

```
6       30445
```

```
5       30445
```

```
4       29231
```

```
2       29043
```

```
Name: count, dtype: int64
```

Extract state code and zip code from address

```
data[['address']]
{"type": "dataframe"}
data['address'].apply(lambda x: x[-8:-6:].strip())
0      OK
1      SD
2      WV
3      MA
4      VA
...
396025  DC
396026  LA
396027  NY
396028  FL
396029  AR
Name: address, Length: 396030, dtype: object

data['state'] = data['address'].apply(lambda x: x[-8:-6:].strip())
data['address'].apply(lambda x: x[-5::].strip())
0      22690
1      05113
2      05113
3      00813
4      11650
...
396025  30723
396026  05113
396027  70466
396028  29597
396029  48052
Name: address, Length: 396030, dtype: object

data['zip_code'] = data['address'].apply(lambda x: x[-5::].strip())
data.head()
{"type": "dataframe", "variable_name": "data"}
```

Distribution of loan status among state code and zip code

```
data['state'].value_counts()
state
AP      14308
AE      14157
AA      13919
```

NJ	7091
WI	7081
LA	7068
NV	7038
AK	7034
MA	7022
VA	7022
VT	7005
NY	7004
MS	7003
TX	7000
SC	6973
ME	6972
AR	6969
OH	6969
GA	6967
ID	6958
IN	6958
KS	6945
WV	6944
RI	6940
MO	6939
IL	6934
WY	6933
NE	6927
HI	6927
IA	6926
FL	6921
AZ	6918
CO	6914
OK	6911
CT	6904
MN	6904
NC	6901
OR	6898
CA	6898
AL	6898
MD	6896
WA	6895
UT	6887
SD	6887
MT	6883
DE	6874
TN	6869
ND	6858
MI	6854
DC	6842
NM	6842
PA	6825

```
NH      6818
KY      6800
Name: count, dtype: int64


data['zip_code'].value_counts()

zip_code
70466    56985
30723    56546
22690    56527
48052    55917
00813    45824
29597    45471
05113    45402
11650    11226
93700    11151
86630    10981
Name: count, dtype: int64

data['zip_code'].nunique()


10

fig, ax = plt.subplots(figsize=(10,10))
sns.countplot(data = data, x = 'state', hue = 'loan_status', palette =
'ocean', ax = ax)
plt.xticks(rotation = 90)
plt.show()
```



1. State codes 'AP', 'AE' and 'AA' are the top 3 states from which loan applications have been received (in same order).
2. For all other states, it's almost similar.
3. Thus, distribution is different across states.

```
fig, ax = plt.subplots(figsize=(10,10))
sns.countplot(data = data, x = 'zip_code', hue = 'loan_status',
palette = 'ocean', ax = ax)
plt.xticks(rotation = 90)
plt.show()
```



1. We can see that the distribution of borrowers w.r.t. their loan status is significantly different as per the zip codes.
2. Zip codes: 05113, 00813, 29597 are having only non-defaulters where as the zip codes: 11650, 86630, 93700 are having only defaulters.

Creation of flags for Pub_rec, Mort_acc and Pub_rec_bankruptcies


```
data['is_pub_rec'] = np.where(data['pub_rec'] > 1.0, 1, 0)
data['is_mort_acc'] = np.where(data['mort_acc'] > 1.0, 1, 0)
data['is_pub_rec_bankruptcies'] =
np.where(data['pub_rec_bankruptcies'] > 1.0, 1, 0)
```



```
data.head()

{"type": "dataframe", "variable_name": "data"}

cols = ['is_pub_rec', 'is_mort_acc', 'is_pub_rec_bankruptcies']

for col in cols:
    sns.countplot(data = data, x = col, fill = 'True', hue =
'loan_status')
    plt.xticks(rotation = 90)
    plt.show()
```





We can see that when either negative records on borrower's public credit profile are present or bankruptcy records are available for borrower, then there are more chances of defaulting.

Data preparation for modeling

```
data.columns
```

```
Index(['loan_amnt', 'term', 'int_rate', 'grade', 'sub_grade',  
      'home_ownership',  
        'annual_inc', 'verification_status', 'issue_d', 'loan_status',  
        'purpose', 'dti', 'earliest_cr_line', 'open_acc', 'pub_rec',  
        'revol_bal', 'total_acc', 'initial_list_status',  
      'application_type',  
        'address', 'emp_title', 'title', 'mort_acc', 'revol_util',  
        'pub_rec_bankruptcies', 'is_defaulter', 'emp_length',  
      'issue_d_month',  
        'issue_d_year', 'earliest_cr_line_month',  
      'earliest_cr_line_year',  
        'issue_d_month_no', 'earliest_cr_line_month_no', 'state',  
      'zip_code',  
        'is_pub_rec', 'is_mort_acc', 'is_pub_rec_bankruptcies'],  
      dtype='object')
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 396030 entries, 0 to 396029  
Data columns (total 38 columns):
```

#	Column	Non-Null Count	Dtype
0	loan_amnt	396030 non-null	float64
1	term	396030 non-null	object
2	int_rate	396030 non-null	float64
3	grade	396030 non-null	object
4	sub_grade	396030 non-null	object
5	home_ownership	396030 non-null	object
6	annual_inc	396030 non-null	float64
7	verification_status	396030 non-null	object
8	issue_d	396030 non-null	object
9	loan_status	396030 non-null	object
10	purpose	396030 non-null	object
11	dti	396030 non-null	float64
12	earliest_cr_line	396030 non-null	object
13	open_acc	396030 non-null	float64
14	pub_rec	396030 non-null	float64
15	revol_bal	396030 non-null	float64
16	total_acc	396030 non-null	float64
17	initial_list_status	396030 non-null	object
18	application_type	396030 non-null	object
19	address	396030 non-null	object

```

20 emp_title          396030 non-null object
21 title              396030 non-null object
22 mort_acc           396030 non-null float64
23 revol_util         396030 non-null float64
24 pub_rec_bankruptcies 396030 non-null float64
25 is_defaulter        396030 non-null float64
26 emp_length         396030 non-null float64
27 issue_d_month       396030 non-null object
28 issue_d_year        396030 non-null object
29 earliest_cr_line_month 396030 non-null object
30 earliest_cr_line_year 396030 non-null object
31 issue_d_month_no    396030 non-null int64
32 earliest_cr_line_month_no 396030 non-null int64
33 state              396030 non-null object
34 zip_code            396030 non-null object
35 is_pub_rec          396030 non-null int64
36 is_mort_acc         396030 non-null int64
37 is_pub_rec_bankruptcies 396030 non-null int64
dtypes: float64(13), int64(5), object(20)
memory usage: 114.8+ MB

```

Dropping extra unnecessary columns

```

data.drop(columns = ['pub_rec', 'mort_acc', 'pub_rec_bankruptcies',
'issue_d_month', 'earliest_cr_line_month', 'issue_d', 'loan_status',
'earliest_cr_line', 'address'], axis = 1, inplace = True)

```

```
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 29 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                             396030 non-null float64
1   term                                  396030 non-null object
2   int_rate                              396030 non-null float64
3   grade                                 396030 non-null object
4   sub_grade                             396030 non-null object
5   home_ownership                        396030 non-null object
6   annual_inc                            396030 non-null float64
7   verification_status                   396030 non-null object
8   purpose                               396030 non-null object
9   dti                                    396030 non-null float64
10  open_acc                              396030 non-null float64
11  revol_bal                             396030 non-null float64
12  total_acc                             396030 non-null float64
13  initial_list_status                   396030 non-null object
14  application_type                       396030 non-null object

```

```

15 emp_title      396030 non-null object
16 title          396030 non-null object
17 revol_util     396030 non-null float64
18 is_defaulter   396030 non-null float64
19 emp_length     396030 non-null float64
20 issue_d_year   396030 non-null object
21 earliest_cr_line_year 396030 non-null object
22 issue_d_month_no 396030 non-null int64
23 earliest_cr_line_month_no 396030 non-null int64
24 state          396030 non-null object
25 zip_code       396030 non-null object
26 is_pub_rec     396030 non-null int64
27 is_mort_acc    396030 non-null int64
28 is_pub_rec_bankruptcies 396030 non-null int64
dtypes: float64(10), int64(5), object(14)
memory usage: 87.6+ MB

```

```

data = data.rename(columns = {'is_pub_rec' : 'pub_rec',
'is_mort_acc' : 'mort_acc', 'is_pub_rec_bankruptcies' :
'pub_rec_bankruptcies'})

```

```
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 29 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                             396030 non-null float64
1   term                                  396030 non-null object
2   int_rate                              396030 non-null float64
3   grade                                 396030 non-null object
4   sub_grade                             396030 non-null object
5   home_ownership                        396030 non-null object
6   annual_inc                            396030 non-null float64
7   verification_status                  396030 non-null object
8   purpose                               396030 non-null object
9   dti                                   396030 non-null float64
10  open_acc                              396030 non-null float64
11  revol_bal                             396030 non-null float64
12  total_acc                             396030 non-null float64
13  initial_list_status                  396030 non-null object
14  application_type                     396030 non-null object
15  emp_title                             396030 non-null object
16  title                                 396030 non-null object
17  revol_util                           396030 non-null float64
18  is_defaulter                         396030 non-null float64
19  emp_length                           396030 non-null float64
20  issue_d_year                         396030 non-null object
21  earliest_cr_line_year                396030 non-null object

```

```
22  issue_d_month_no          396030 non-null  int64
23  earliest_cr_line_month_no 396030 non-null  int64
24  state                     396030 non-null  object
25  zip_code                   396030 non-null  object
26  pub_rec                    396030 non-null  int64
27  mort_acc                   396030 non-null  int64
28  pub_rec_bankruptcies      396030 non-null  int64
dtypes: float64(10), int64(5), object(14)
memory usage: 87.6+ MB
```

Converting necessary columns to category.

```
for col in data.columns:
    if data[col].dtype == 'object':
        data[col] = data[col].astype('category')

for col in data.columns:
    if data[col].dtype == 'category':
        print(col)

term
grade
sub_grade
home_ownership
verification_status
purpose
initial_list_status
application_type
emp_title
title
issue_d_year
earliest_cr_line_year
state
zip_code

for col in data.columns:
    if data[col].dtype != 'category':
        print(col)

loan_amnt
int_rate
annual_inc
dti
open_acc
revol_bal
total_acc
revol_util
is_defaulter
emp_length
issue_d_month_no
```

```

earliest_cr_line_month_no
pub_rec
mort_acc
pub_rec_bankruptcies

columns_to_cat =
['emp_length', 'issue_d_month_no', 'earliest_cr_line_month_no', 'pub_rec'
, 'mort_acc', 'pub_rec_bankruptcies']

for col in columns_to_cat:
    data[col] = data[col].astype('category')

data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 29 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                            396030 non-null float64
1   term                                396030 non-null category
2   int_rate                            396030 non-null float64
3   grade                               396030 non-null category
4   sub_grade                           396030 non-null category
5   home_ownership                       396030 non-null category
6   annual_inc                           396030 non-null float64
7   verification_status                 396030 non-null category
8   purpose                              396030 non-null category
9   dti                                  396030 non-null float64
10  open_acc                             396030 non-null float64
11  revol_bal                            396030 non-null float64
12  total_acc                            396030 non-null float64
13  initial_list_status                 396030 non-null category
14  application_type                    396030 non-null category
15  emp_title                           396030 non-null category
16  title                               396030 non-null category
17  revol_util                           396030 non-null float64
18  is_defaulter                         396030 non-null float64
19  emp_length                           396030 non-null category
20  issue_d_year                         396030 non-null category
21  earliest_cr_line_year               396030 non-null category
22  issue_d_month_no                    396030 non-null category
23  earliest_cr_line_month_no           396030 non-null category
24  state                               396030 non-null category
25  zip_code                            396030 non-null category
26  pub_rec                             396030 non-null category
27  mort_acc                            396030 non-null category
28  pub_rec_bankruptcies                396030 non-null category
dtypes: category(20), float64(9)
memory usage: 43.7 MB

```

```

target_encoding_cols = ['term', 'grade', 'home_ownership',
                        'verification_status', 'initial_list_status', 'application_type',
                        'sub_grade', 'purpose', 'emp_title', 'title', 'state',
                        'issue_d_month_no', 'earliest_cr_line_month_no', 'issue_d_year',
                        'earliest_cr_line_year', 'zip_code',
                        'pub_rec', 'mort_acc', 'pub_rec_bankruptcies', 'emp_length']

for col in data.columns:
    if data[col].dtype == 'category' and col not in
target_encoding_cols:
        print(col)

for col in data.columns:
    if data[col].dtype != 'category' and col in target_encoding_cols:
        print(col)

```

Target Encoding

We will apply target encoding for all the category columns as one hot encoding for columns with 2 categories might add multi-collinearity. So, it will replace category with mean value of target column i.e. is_defaulter with that category value.

```

target_encoder = ce.TargetEncoder()

data[['term', 'grade', 'home_ownership', 'verification_status',
      'initial_list_status', 'application_type', 'sub_grade', 'purpose',
      'emp_title', 'title', 'state', 'issue_d_month_no',
      'earliest_cr_line_month_no', 'issue_d_year', 'earliest_cr_line_year',
      'zip_code', 'pub_rec', 'mort_acc', 'pub_rec_bankruptcies',
      'emp_length']] = target_encoder.fit_transform(data[['term', 'grade',
      'home_ownership', 'verification_status', 'initial_list_status',
      'application_type', 'sub_grade', 'purpose', 'emp_title', 'title',
      'state', 'issue_d_month_no', 'earliest_cr_line_month_no',
      'issue_d_year', 'earliest_cr_line_year', 'zip_code',
      'pub_rec', 'mort_acc', 'pub_rec_bankruptcies', 'emp_length']],
data['is_defaulter'])

data

{"type": "dataframe", "variable_name": "data"}

```

#Model Building

Split training and testing data

```

data['is_defaulter']

0      0.0
1      0.0
2      0.0

```

```

3          0.0
4          1.0
...
396025     0.0
396026     0.0
396027     0.0
396028     0.0
396029     0.0
Name: is_defaulter, Length: 396030, dtype: float64

X = data.drop('is_defaulter', axis = 1)
y = data['is_defaulter']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.3, random_state = 42)

X_train.shape
(277221, 28)

X_test.shape
(118809, 28)

y_train.shape
(277221,)

y_test.shape
(118809,)

X_train.head()
{"type": "dataframe", "variable_name": "X_train"}

y_train.head()
3412      0.0
134032    0.0
19526     0.0
61015     0.0
2896      0.0
Name: is_defaulter, dtype: float64

X_train.info()
<class 'pandas.core.frame.DataFrame'>
Index: 277221 entries, 3412 to 121958
Data columns (total 28 columns):
#   Column              Non-Null Count  Dtype
---  -
0   loan_amnt           277221 non-null float64

```


1	term	277221	non-null	float64
2	int_rate	277221	non-null	float64
3	grade	277221	non-null	float64
4	sub_grade	277221	non-null	float64
5	home_ownership	277221	non-null	float64
6	annual_inc	277221	non-null	float64
7	verification_status	277221	non-null	float64
8	purpose	277221	non-null	float64
9	dti	277221	non-null	float64
10	open_acc	277221	non-null	float64
11	revol_bal	277221	non-null	float64
12	total_acc	277221	non-null	float64
13	initial_list_status	277221	non-null	float64
14	application_type	277221	non-null	float64
15	emp_title	277221	non-null	float64
16	title	277221	non-null	float64
17	revol_util	277221	non-null	float64
18	emp_length	277221	non-null	float64
19	issue_d_year	277221	non-null	float64
20	earliest_cr_line_year	277221	non-null	float64
21	issue_d_month_no	277221	non-null	float64
22	earliest_cr_line_month_no	277221	non-null	float64
23	state	277221	non-null	float64
24	zip_code	277221	non-null	float64
25	pub_rec	277221	non-null	float64
26	mort_acc	277221	non-null	float64
27	pub_rec_bankruptcies	277221	non-null	float64

dtypes: float64(28)

memory usage: 61.3 MB

y_train.info()

```
<class 'pandas.core.series.Series'>
Index: 277221 entries, 3412 to 121958
Series name: is_defaulter
Non-Null Count  Dtype
-----
277221 non-null  float64
dtypes: float64(1)
memory usage: 4.2 MB
```

y_train

3412	0.0
134032	0.0
19526	0.0
61015	0.0
2896	0.0
...	
259178	0.0

```
365838    1.0
131932    0.0
146867    0.0
121958    1.0
Name: is_defaulter, Length: 277221, dtype: float64
```

Feature Scaling

Using Standard Scaler as scaling method and note that fit_transform should be done only on training data.

```
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

Logistic Regression using sklearn

```
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
LogisticRegression(max_iter=1000)
pred_y_train = model.predict(X_train)
pred_y_test = model.predict(X_test)
print('Accuracy of Logistic Regression Classifier on test set:
{:.3f}'.format(model.score(X_test, y_test)))
Accuracy of Logistic Regression Classifier on test set: 0.920
```

Model coef and intercept

```
coefs = model.coef_.reshape(-1,1)
lst = list(X.columns)

#w1, w2...w27
coefs

array([[ 1.38398980e-01],
       [ 2.29202670e-01],
       [ 1.77839868e-01],
       [-2.33215218e-02],
       [ 2.11443813e-01],
       [ 1.46470923e-01],
       [-6.77768341e-03],
       [ 3.35492924e-02],
       [-1.97819527e-01],
       [ 1.47518410e-01],
       [ 1.73290930e-01],
```

```
[-7.15317429e-02],  
[-7.00802116e-02],  
[-1.17981489e-01],  
[ 4.84440015e-03],  
[ 1.31665904e+00],  
[ 4.92844290e-01],  
[ 2.04307342e-01],  
[ 2.57313199e-02],  
[ 5.44637475e-02],  
[-4.45658436e-03],  
[ 2.77710460e-02],  
[-4.90437871e-03],  
[ 3.15780681e-02],  
[ 9.28179390e+00],  
[-1.16193516e-02],  
[ 8.29470126e-02],  
[-4.73083570e-03]])
```

```
for i in range(len(lst)):  
    print('Column: ', lst[i], ' Coef: ', coefs[i])
```

```
Column: loan_amnt Coef: [0.13839898]  
Column: term Coef: [0.22920267]  
Column: int_rate Coef: [0.17783987]  
Column: grade Coef: [-0.02332152]  
Column: sub_grade Coef: [0.21144381]  
Column: home_ownership Coef: [0.14647092]  
Column: annual_inc Coef: [-0.00677768]  
Column: verification_status Coef: [0.03354929]  
Column: purpose Coef: [-0.19781953]  
Column: dti Coef: [0.14751841]  
Column: open_acc Coef: [0.17329093]  
Column: revol_bal Coef: [-0.07153174]  
Column: total_acc Coef: [-0.07008021]  
Column: initial_list_status Coef: [-0.11798149]  
Column: application_type Coef: [0.0048444]  
Column: emp_title Coef: [1.31665904]  
Column: title Coef: [0.49284429]  
Column: revol_util Coef: [0.20430734]  
Column: emp_length Coef: [0.02573132]  
Column: issue_d_year Coef: [0.05446375]  
Column: earliest_cr_line_year Coef: [-0.00445658]  
Column: issue_d_month_no Coef: [0.02777105]  
Column: earliest_cr_line_month_no Coef: [-0.00490438]  
Column: state Coef: [0.03157807]  
Column: zip_code Coef: [9.2817939]  
Column: pub_rec Coef: [-0.01161935]  
Column: mort_acc Coef: [0.08294701]  
Column: pub_rec_bankruptcies Coef: [-0.00473084]
```

```

coef_df = pd.DataFrame()
coef_df['column_name'] = lst
coef_df['coef_value'] = coefs
coef_df.sort_values(by = 'coef_value', ascending = False)

{"summary":{"\n  \"name\": \"coef_df\",\n  \"rows\": 28,\n  \"fields\": [\n    {\n      \"column\": \"column_name\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 28,\n        \"samples\": [\n          \"home_ownership\",\n          \"revol_bal\",\n          \"dti\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"coef_value\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 1.754294768944616,\n        \"min\": -0.1978195274163468,\n        \"max\": 9.28179390488985,\n        \"num_unique_values\": 28,\n        \"samples\": [\n          0.14647092257489297,\n          -\n          0.07153174292788539,\n          0.14751841013769154\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ],\n  \"type\": \"dataframe\"}

```

We can see that zip_code and emp_title are having top 2 highest weights, which is quite surprising.

```

#w0
model.intercept_
array([-2.11207663])
model.score(X_test, y_test)
0.9204689880396266

```

Metric evaluation

Confusion Matrix

```


conf_matrix = confusion_matrix(y_test, pred_y_test)
conf_matrix

array([[92694, 2745],
       [ 6704, 16666]])

fig, ax = plt.subplots(figsize= (5,5))
ConfusionMatrixDisplay(conf_matrix).plot(ax = ax)

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7bd14fadbd30>

```



Accuracy

```
accuracy = np.diag(conf_matrix).sum() / conf_matrix.sum()  
accuracy
```

```
0.9204689880396266
```

```
accuracy_score(y_test, pred_y_test)
```

```
0.9204689880396266
```

Precision

```
precision_score(y_test, pred_y_test)
```

```
0.8585853382102931
```

Recall

```
recall_score(y_test, pred_y_test)
```

```
0.7131364997860505
```

Precision is higher than recall which means that FN is higher than FP.

F1 Score

Testing F1-score

```
f1_score(y_test, pred_y_test)
0.7791309226058298
```

Training F1-score

```
f1_score(y_train, pred_y_train)
0.7793166011490779
```

Classification Report

```
from sklearn.metrics import classification_report
print(classification_report(y_test, pred_y_test))
```

	precision	recall	f1-score	support
0.0	0.93	0.97	0.95	95439
1.0	0.86	0.71	0.78	23370
accuracy			0.92	118809
macro avg	0.90	0.84	0.87	118809
weighted avg	0.92	0.92	0.92	118809

AUC-ROC Curve


```
probability = model.predict_proba(X_test)
probability
array([[3.22092916e-01, 6.77907084e-01],
       [9.99385798e-01, 6.14202397e-04],
       [9.32259580e-01, 6.77404198e-02],
       ...,
       [9.63454985e-01, 3.65450145e-02],
       [9.99344302e-01, 6.55697673e-04],
       [9.99937408e-01, 6.25923859e-05]])

probabilities = probability[:, 1]
probabilities
array([6.77907084e-01, 6.14202397e-04, 6.77404198e-02, ...,
       3.65450145e-02, 6.55697673e-04, 6.25923859e-05])

fpr, tpr, thr = roc_curve(y_test, probabilities)
plt.plot(fpr, tpr)

#random model
plt.plot(fpr, fpr, '--', color = 'red')
```

```
plt.title('ROC Curve')
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.show()
```




```
roc_auc_score(y_test, probabilities)
0.9597543106693195
```

PR Curve

```
precision, recall, thr = precision_recall_curve(y_test, probabilities)
plt.plot(recall, precision)

plt.xlabel('recall')
plt.ylabel('precision')
plt.title('PR Curve')
plt.show()
```



```
auc(recall, precision)
0.8881669306885851

def precision_recall_curve_plot(y_test, pred_proba_c1):
    precisions, recalls, thresholds = precision_recall_curve(y_test,
pred_proba_c1)


    threshold_boundary = thresholds.shape[0]
    # plot precision
    plt.plot(thresholds, precisions[0:threshold_boundary],
linestyle='--', label='precision')
    # plot recall
    plt.plot(thresholds, recalls[0:threshold_boundary],
label='recalls')

    start, end = plt.xlim()
    plt.xticks(np.round(np.arange(start, end, 0.1), 2))

    plt.xlabel('Threshold Value'); plt.ylabel('Precision and Recall
Value')
    plt.legend(); plt.grid()
    plt.show()
```



```
precision_recall_curve_plot(y_test, model.predict_proba(X_test)[: ,1])
```



Threshold value here comes out to be almost 0.35

Multicollinearity check using Variance Inflation Factor (VIF)

-

Multicollinearity occurs when two or more independent variables are highly correlated with one another in a regression model. Multicollinearity can be a problem in a regression model because we would not be able to distinguish between the individual effects of the independent variables on the dependent variable.

Multicollinearity can be detected via various methods. One such method is Variance Inflation Factor aka VIF. In VIF method, we pick each independent feature and regress it against all of the other independent features. VIF score of an independent variable represents how well the variable is explained by other independent variables.

$VIF = 1/(1-R^2)$

```
def calc_vif(X):  
    # Calculating the VIF  
    vif = pd.DataFrame()
```

```

vif['Feature'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in
range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by='VIF', ascending = False)
return vif

calc_vif(X)

{"summary":{"\n  \"name\": \"calc_vif(X)\",\n  \"rows\": 28,\n  \"fields\": [\n    {\n      \"column\": \"Feature\",
\"properties\": {\n        \"dtype\": \"string\",
\"num_unique_values\": 28,\n        \"samples\": [\n          \"sub_grade\",
          \"loan_amnt\",
          \"int_rate\"
        ],\n        \"semantic_type\": \"\",
        \"description\": \"\"
      },\n      {\n        \"column\": \"VIF\",
        \"properties\": {\n          \"dtype\": \"number\",
          \"std\": 1445.469758322807,\n          \"min\": 1.81,\n          \"max\": 5966.65,\n          \"num_unique_values\": 28,\n          \"samples\": [\n            187.51,\n            7.17,\n            212.17
          ],\n          \"semantic_type\": \"\",
          \"description\": \"\"
        }
      ]
    },\n    {\n      \"column\": \"VIF\",
      \"properties\": {\n        \"dtype\": \"number\",
        \"std\": 1445.469758322807,\n        \"min\": 1.81,\n        \"max\": 5966.65,\n        \"num_unique_values\": 28,\n        \"samples\": [\n          187.51,\n          7.17,\n          212.17
        ],\n        \"semantic_type\": \"\",
        \"description\": \"\"
      }
    ]
  ],\n  \"type\": \"dataframe\"}

X.drop(columns=['pub_rec_bankruptcies'], axis=1, inplace=True)
calc_vif(X)[:5]

{"summary":{"\n  \"name\": \"calc_vif(X)[:5]\",\n  \"rows\": 5,\n  \"fields\": [\n    {\n      \"column\": \"Feature\",
\"properties\": {\n        \"dtype\": \"string\",
\"num_unique_values\": 5,\n        \"samples\": [\n          \"initial_list_status\",
          \"application_type\",
          \"emp_length\",
          \"semantic_type\": \"\",
          \"description\": \"\"
        },\n      {\n        \"column\": \"VIF\",
        \"properties\": {\n          \"dtype\": \"number\",
          \"std\": 879.0135631604327,\n          \"min\": 1475.18,\n          \"max\": 3703.96,\n          \"num_unique_values\": 5,\n          \"samples\": [\n            2533.0,\n            1475.18,\n            2006.58
          ],\n          \"semantic_type\": \"\",
          \"description\": \"\"
        }
      ]
    },\n    {\n      \"column\": \"VIF\",
      \"properties\": {\n        \"dtype\": \"number\",
        \"std\": 879.0135631604327,\n        \"min\": 1475.18,\n        \"max\": 3703.96,\n        \"num_unique_values\": 5,\n        \"samples\": [\n          2533.0,\n          1475.18,\n          2006.58
        ],\n        \"semantic_type\": \"\",
        \"description\": \"\"
      }
    ]
  ],\n  \"type\": \"dataframe\"}

X.drop(columns=['earliest_cr_line_month_no'], axis=1, inplace=True)
calc_vif(X)[:5]

{"summary":{"\n  \"name\": \"calc_vif(X)[:5]\",\n  \"rows\": 5,\n  \"fields\": [\n    {\n      \"column\": \"Feature\",
\"properties\": {\n        \"dtype\": \"string\",
\"num_unique_values\": 5,\n        \"samples\": [\n          \"emp_length\",
          \"pub_rec\",
          \"state\"
        ],\n        \"semantic_type\": \"\",
        \"description\": \"\"
      },\n      {\n        \"column\": \"VIF\",
        \"properties\": {\n          \"dtype\": \"number\",
          \"std\": 879.0135631604327,\n          \"min\": 1475.18,\n          \"max\": 3703.96,\n          \"num_unique_values\": 5,\n          \"samples\": [\n            2533.0,\n            1475.18,\n            2006.58
          ],\n          \"semantic_type\": \"\",
          \"description\": \"\"
        }
      ]
    },\n    {\n      \"column\": \"VIF\",
      \"properties\": {\n        \"dtype\": \"number\",
        \"std\": 879.0135631604327,\n        \"min\": 1475.18,\n        \"max\": 3703.96,\n        \"num_unique_values\": 5,\n        \"samples\": [\n          2533.0,\n          1475.18,\n          2006.58
        ],\n        \"semantic_type\": \"\",
        \"description\": \"\"
      }
    ]
  ],\n  \"type\": \"dataframe\"}

```

```

\"std\": 400.4819925165176,\n      \"min\": 1250.02,\n\"max\": 2267.53,\n      \"num_unique_values\": 5,\n\"samples\": [\n      1849.24,\n      1250.02,\n      1636.76\n      ],\n      \"semantic_type\": \"\",\n\"description\": \"\"\n      }\n      }\n      ]\n      }\", \"type\": \"dataframe\"}

```

```

X.drop(columns=['initial_list_status'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{\"summary\": \"{\\n  \\\"name\\\": \\\"calc_vif(X)[:5]\\\",\\n  \\\"rows\\\": 5,\\n
\\\"fields\\\": [\\n    {\\n      \\\"column\\\": \\\"Feature\\\",\\n
\\\"properties\\\": {\\n        \\\"dtype\\\": \\\"string\\\",\\n
\\\"num_unique_values\\\": 5,\\n        \\\"samples\\\": [\\n
\\\"state\\\",\\n        \\\"issue_d_month_no\\\",\\n
\\\"application_type\\\",\\n        \\\"semantic_type\\\": \\\"\\\",\\n
\\\"description\\\": \\\"\\\"\\n      }\\n    },\\n    {\\n      \\\"column\\\":
\\\"VIF\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"number\\\",\\n
\\\"std\\\": 490.04208693743846,\\n        \\\"min\\\": 409.24,\\n
\\\"max\\\": 1674.54,\\n        \\\"num_unique_values\\\": 5,\\n
\\\"samples\\\": [\\n        1513.27,\\n        409.24,\\n
        1313.33\\n      ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n
\\\"description\\\": \\\"\\\"\\n      }\\n    }\\n  ]\\n}\", \"type\": \"dataframe\"}

```

```

X.drop(columns=['emp_length'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{\"summary\": \"{\\n  \\\"name\\\": \\\"calc_vif(X)[:5]\\\",\\n  \\\"rows\\\": 5,\\n
\\\"fields\\\": [\\n    {\\n      \\\"column\\\": \\\"Feature\\\",\\n
\\\"properties\\\": {\\n        \\\"dtype\\\": \\\"string\\\",\\n
\\\"num_unique_values\\\": 5,\\n        \\\"samples\\\": [\\n
\\\"application_type\\\",\\n        \\\"int_rate\\\",\\n        \\\"pub_rec\\\"\\n
\\n      ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n
\\\"description\\\": \\\"\\\"\\n      }\\n    },\\n    {\\n      \\\"column\\\":
\\\"VIF\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"number\\\",\\n
\\\"std\\\": 499.38102283727204,\\n        \\\"min\\\": 210.15,\\n
\\\"max\\\": 1319.1,\\n        \\\"num_unique_values\\\": 5,\\n
\\\"samples\\\": [\\n        1187.17,\\n        210.15,\\n
        1082.78\\n      ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n
\\\"description\\\": \\\"\\\"\\n      }\\n    }\\n  ]\\n}\", \"type\": \"dataframe\"}

```

```

X.drop(columns=['state'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{\"summary\": \"{\\n  \\\"name\\\": \\\"calc_vif(X)[:5]\\\",\\n  \\\"rows\\\": 5,\\n
\\\"fields\\\": [\\n    {\\n      \\\"column\\\": \\\"Feature\\\",\\n
\\\"properties\\\": {\\n        \\\"dtype\\\": \\\"string\\\",\\n
\\\"num_unique_values\\\": 5,\\n        \\\"samples\\\": [\\n
\\\"pub_rec\\\",\\n        \\\"sub_grade\\\",\\n
\\\"issue_d_month_no\\\",\\n        \\\"semantic_type\\\": \\\"\\\",\\n
\\\"description\\\": \\\"\\\"\\n      }\\n    },\\n    {\\n      \\\"column\\\":
\\\"VIF\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"number\\\",\\n

```

```

{"std": 377.087264210819, "min": 185.42, "max": 961.89, "num_unique_values": 5, "samples": [904.44, 185.42, 387.45], "semantic_type": "", "description": ""}
{"type": "dataframe"}

```

```

X.drop(columns=['application_type'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{"summary": {"name": "calc_vif(X)[:5]", "rows": 5, "fields": [{"column": "Feature", "properties": {"dtype": "string", "num_unique_values": 5, "samples": ["issue_d_month_no", "mort_acc", "int_rate"], "semantic_type": "", "description": ""}], "VIF": {"properties": {"dtype": "number", "std": 167.35722467225608, "min": 123.39, "max": 539.96, "num_unique_values": 5, "samples": [352.78, 123.39, 206.63], "semantic_type": "", "description": ""}], "type": "dataframe"}

```

```

X.drop(columns=['pub_rec'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{"summary": {"name": "calc_vif(X)[:5]", "rows": 5, "fields": [{"column": "Feature", "properties": {"dtype": "string", "num_unique_values": 5, "samples": ["int_rate", "grade", "sub_grade"], "semantic_type": "", "description": ""}], "VIF": {"properties": {"dtype": "number", "std": 50.7330841759103, "min": 105.41, "max": 221.72, "num_unique_values": 5, "samples": [196.92, 105.41, 178.17], "semantic_type": "", "description": ""}], "type": "dataframe"}

```

```

X.drop(columns=['issue_d_month_no'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{"summary": {"name": "calc_vif(X)[:5]", "rows": 5, "fields": [{"column": "Feature", "properties": {"dtype": "string", "num_unique_values": 5, "samples": ["sub_grade", "earliest_cr_line_year", "mort_acc"], "semantic_type": "", "description": ""}], "VIF": {"properties": {"dtype": "number",

```

```

{"std\\": 42.96233664501967,\\n          \\\"min\\\": 87.23,\\n          \\\"max\\\": 182.73,\\n          \\\"num_unique_values\\\": 5,\\n          \\\"samples\\\": [\\n 169.56,\\n          87.23,\\n          105.46\\n          ],\\n          \\\"semantic_type\\\": \\\"\\\",\\n          \\\"description\\\": \\\"\\\"\\n          }\\n    }\\n  ]\\n}\\",\"type\":\"dataframe\"}

```

```

X.drop(columns=['int_rate'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{"summary": "{\\n  \\\"name\\\": \\\"calc_vif(X)[:5]\\\",\\n  \\\"rows\\\": 5,\\n  \\\"fields\\\": [\\n    {\\n      \\\"column\\\": \\\"Feature\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"string\\\",\\n        \\\"num_unique_values\\\": 5,\\n        \\\"samples\\\": [\\n          \\\"mort_acc\\\",\\n          \\\"purpose\\\",\\n          \\\"sub_grade\\\"\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n        \\\"description\\\": \\\"\\\"\\n      }\\n    },\\n    {\\n      \\\"column\\\": \\\"VIF\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"number\\\",\\n        \\\"std\\\": 12.236969396055542,\\n        \\\"min\\\": 80.92,\\n        \\\"max\\\": 105.26,\\n        \\\"num_unique_values\\\": 5,\\n        \\\"samples\\\": [\\n          104.74,\\n          80.92,\\n          104.27\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n        \\\"description\\\": \\\"\\\"\\n      }\\n    }\\n  ]\\n}\\",\"type\":\"dataframe\"}

```

```

X.drop(columns=['grade'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{"summary": "{\\n  \\\"name\\\": \\\"calc_vif(X)[:5]\\\",\\n  \\\"rows\\\": 5,\\n  \\\"fields\\\": [\\n    {\\n      \\\"column\\\": \\\"Feature\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"string\\\",\\n        \\\"num_unique_values\\\": 5,\\n        \\\"samples\\\": [\\n          \\\"earliest_cr_line_year\\\",\\n          \\\"title\\\",\\n          \\\"purpose\\\"\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n        \\\"description\\\": \\\"\\\"\\n      }\\n    },\\n    {\\n      \\\"column\\\": \\\"VIF\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"number\\\",\\n        \\\"std\\\": 21.565842900290267,\\n        \\\"min\\\": 46.07,\\n        \\\"max\\\": 104.73,\\n        \\\"num_unique_values\\\": 5,\\n        \\\"samples\\\": [\\n          84.09,\\n          46.07,\\n          80.91\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n        \\\"description\\\": \\\"\\\"\\n      }\\n    }\\n  ]\\n}\\",\"type\":\"dataframe\"}

```

```

X.drop(columns=['mort_acc'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{"summary": "{\\n  \\\"name\\\": \\\"calc_vif(X)[:5]\\\",\\n  \\\"rows\\\": 5,\\n  \\\"fields\\\": [\\n    {\\n      \\\"column\\\": \\\"Feature\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"string\\\",\\n        \\\"num_unique_values\\\": 5,\\n        \\\"samples\\\": [\\n          \\\"earliest_cr_line_year\\\",\\n          \\\"verification_status\\\",\\n          \\\"home_ownership\\\"\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n        \\\"description\\\": \\\"\\\"\\n      }\\n    },\\n    {\\n      \\\"column\\\": \\\"VIF\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"number\\\",\\n

```

```

\"std\": 17.590588392660436,\n      \"min\": 36.52,\n      \"max\": 77.25,\n      \"num_unique_values\": 5,\n      \"samples\": [\n        72.56,\n        36.52,\n        50.26\n      ],\n      \"semantic_type\": \"\",\n      \"description\": \"\"\n    }\n  ],\n  \"type\": \"dataframe\"}

```

```

X.drop(columns=['purpose'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{"summary": "{\n  \"name\": \"calc_vif(X)[:5]\",\n  \"rows\": 5,\n  \"fields\": [\n    {\n      \"column\": \"Feature\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 5,\n        \"samples\": [\n          \"home_ownership\",\n          \"issue_d_year\",\n          \"verification_status\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"VIF\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 16.727715624077305,\n        \"min\": 23.3,\n        \"max\": 67.31,\n        \"num_unique_values\": 5,\n        \"samples\": [\n          47.5,\n          23.3,\n          35.6\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ],\n  \"type\": \"dataframe\"}

```

```

X.drop(columns=['earliest_cr_line_year'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{"summary": "{\n  \"name\": \"calc_vif(X)[:5]\",\n  \"rows\": 5,\n  \"fields\": [\n    {\n      \"column\": \"Feature\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 5,\n        \"samples\": [\n          \"verification_status\",\n          \"emp_title\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"VIF\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 7.828088527859148,\n        \"min\": 16.75,\n        \"max\": 33.7,\n        \"num_unique_values\": 5,\n        \"samples\": [\n          33.68,\n          16.75,\n          33.05\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ],\n  \"type\": \"dataframe\"}

```

```

X.drop(columns=['home_ownership'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{"summary": "{\n  \"name\": \"calc_vif(X)[:5]\",\n  \"rows\": 5,\n  \"fields\": [\n    {\n      \"column\": \"Feature\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 5,\n        \"samples\": [\n          \"verification_status\",\n          \"open_acc\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"VIF\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 7.828088527859148,\n        \"min\": 16.75,\n        \"max\": 33.7,\n        \"num_unique_values\": 5,\n        \"samples\": [\n          33.68,\n          16.75,\n          33.05\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ],\n  \"type\": \"dataframe\"}

```

```

{"std\\": 7.455894983166005,\\n          \\\"min\\\": 13.83,\\n          \\\"max\\\": 30.42,\\n          \\\"num_unique_values\\\": 5,\\n          \\\"samples\\\": [\\n 29.09,\\n          13.83,\\n          22.16\\n          ],\\n          \\\"semantic_type\\\": \\\"\\\",\\n          \\\"description\\\": \\\"\\\"\\n          }\\n    }\\n  ]\\n}\\",\"type\":\"dataframe\"}

```

```

X.drop(columns=['title'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{"summary": "{\\n  \\\"name\\\": \\\"calc_vif(X)[:5]\\\",\\n  \\\"rows\\\": 5,\\n  \\\"fields\\\": [\\n    {\\n      \\\"column\\\": \\\"Feature\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"string\\\",\\n        \\\"num_unique_values\\\": 5,\\n        \\\"samples\\\": [\\n          \\\"issue_d_year\\\",\\n          \\\"term\\\",\\n          \\\"emp_title\\\"\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n        \\\"description\\\": \\\"\\\"\\n      }\\n    },\\n    {\\n      \\\"column\\\": \\\"VIF\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"number\\\",\\n        \\\"std\\\": 5.671773972929458,\\n        \\\"min\\\": 13.05,\\n        \\\"max\\\": 26.91,\\n        \\\"num_unique_values\\\": 5,\\n        \\\"samples\\\": [\\n          18.17,\\n          13.05,\\n          15.03\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n        \\\"description\\\": \\\"\\\"\\n      }\\n    }\\n  ]\\n}\\",\"type\":\"dataframe\"}

```

```

X.drop(columns=['verification_status'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{"summary": "{\\n  \\\"name\\\": \\\"calc_vif(X)[:5]\\\",\\n  \\\"rows\\\": 5,\\n  \\\"fields\\\": [\\n    {\\n      \\\"column\\\": \\\"Feature\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"string\\\",\\n        \\\"num_unique_values\\\": 5,\\n        \\\"samples\\\": [\\n          \\\"open_acc\\\",\\n          \\\"total_acc\\\",\\n          \\\"emp_title\\\"\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n        \\\"description\\\": \\\"\\\"\\n      }\\n    },\\n    {\\n      \\\"column\\\": \\\"VIF\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"number\\\",\\n        \\\"std\\\": 1.821930295044242,\\n        \\\"min\\\": 11.78,\\n        \\\"max\\\": 16.58,\\n        \\\"num_unique_values\\\": 5,\\n        \\\"samples\\\": [\\n          13.71,\\n          11.78,\\n          13.37\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n        \\\"description\\\": \\\"\\\"\\n      }\\n    }\\n  ]\\n}\\",\"type\":\"dataframe\"}

```

```

X.drop(columns=['issue_d_year'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{"summary": "{\\n  \\\"name\\\": \\\"calc_vif(X)[:5]\\\",\\n  \\\"rows\\\": 5,\\n  \\\"fields\\\": [\\n    {\\n      \\\"column\\\": \\\"Feature\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"string\\\",\\n        \\\"num_unique_values\\\": 5,\\n        \\\"samples\\\": [\\n          \\\"term\\\",\\n          \\\"annual_inc\\\",\\n          \\\"total_acc\\\"\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n        \\\"description\\\": \\\"\\\"\\n      }\\n    },\\n    {\\n      \\\"column\\\": \\\"VIF\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"number\\\",\\n        \\\"std\\\": 1.9949887217726319,\\n

```

```

n          \ "min\ ": 8.17,\n          \ "max\ ": 13.51,\n
\ "num_unique_values\ ": 5,\n          \ "samples\ ": [\n          12.13,\n
8.17,\n          11.76\n          ],\n          \ "semantic_type\ ": \ "\",\n
\ "description\ ": \ "\n          }\n          }\n          ]\n          }", "type": "dataframe"}

```

```

X.drop(columns=['open_acc'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{"summary": "{\n  \ "name\ ": \ "calc_vif(X)[:5]\ ",\n  \ "rows\ ": 5,\n
\ "fields\ ": [\n    {\n      \ "column\ ": \ "Feature\ ",\n
\ "properties\ ": {\n        \ "dtype\ ": \ "string\ ",\n
\ "num_unique_values\ ": 5,\n        \ "samples\ ": [\n
\ "emp_title\ ",\n        \ "revol_util\ ",\n        \ "annual_inc\ "\n
],\n        \ "semantic_type\ ": \ "\",\n        \ "description\ ": \ "\n
}\n      },\n      {\n        \ "column\ ": \ "VIF\ ",\n        \ "properties\ ": {\n
n          \ "dtype\ ": \ "number\ ",\n          \ "std\ ": 2.129143020090478,\n
\ "min\ ": 7.3,\n          \ "max\ ": 12.13,\n          \ "num_unique_values\ ":
5,\n          \ "samples\ ": [\n          10.47,\n          7.3,\n
8.06\n          ],\n          \ "semantic_type\ ": \ "\",\n
\ "description\ ": \ "\n          }\n          }\n          ]\n          }", "type": "dataframe"}

```

```

X.drop(columns=['term'], axis=1, inplace=True)
calc_vif(X)[:5]

```

```

{"summary": "{\n  \ "name\ ": \ "calc_vif(X)[:5]\ ",\n  \ "rows\ ": 5,\n
\ "fields\ ": [\n    {\n      \ "column\ ": \ "Feature\ ",\n
\ "properties\ ": {\n        \ "dtype\ ": \ "string\ ",\n
\ "num_unique_values\ ": 5,\n        \ "samples\ ": [\n
\ "annual_inc\ ",\n        \ "dti\ ",\n        \ "total_acc\ "\n
n          ],\n        \ "semantic_type\ ": \ "\",\n
\ "description\ ": \ "\n          }\n          },\n      {\n        \ "column\ ":
\ "VIF\ ",\n        \ "properties\ ": {\n          \ "dtype\ ": \ "number\ ",\n
\ "std\ ": 1.0531476629608976,\n          \ "min\ ": 7.1,\n          \ "max\ ":
9.66,\n          \ "num_unique_values\ ": 5,\n          \ "samples\ ": [\n
7.98,\n          7.1,\n          7.32\n          ],\n          \ "semantic_type\ ": \ "\",\n
\ "description\ ": \ "\n          }\n          }\n          ]\n          }", "type": "dataframe"}

```

```

X.head()

```

```

{"type": "dataframe", "variable_name": "X"}

```

```

X = scaler.fit_transform(X)

```

```

kfold = KFold(n_splits=5)
accuracy = np.mean(cross_val_score(model, X, y, cv=kfold,
scoring='accuracy', n_jobs=-1))
print("Cross Validation accuracy: {:.3f}".format(accuracy))

```

```

Cross Validation accuracy: 0.919

```


SMOTE

Oversampling by creating synthetic samples for minority class which is class 1 here i.e. defaulters to make the no of samples for class 1 same as class 0

```
smt = SMOTE()

print('Before SMOTE')
y_train.value_counts()

Before SMOTE
is_defaulter
0.0    222918
1.0     54303
Name: count, dtype: int64

X_sm, y_sm = smt.fit_resample(X_train, y_train)

print('After SMOTE')
y_sm.value_counts()

After SMOTE
is_defaulter
0.0    222918
1.0    222918
Name: count, dtype: int64

model_smote= LogisticRegression(max_iter=1000)

def training(model, X_train, X_test, y_train, y_test):
    model.fit(X_train, y_train)

    train_y_pred = model.predict(X_train)
    test_y_pred = model.predict(X_test)

    train_score = f1_score(y_train, train_y_pred)
    test_score = f1_score(y_test, test_y_pred)

    print(classification_report(y_test, test_y_pred))

    return train_score, test_score

f1_train, f1_test = training(model_smote, X_sm, X_test, y_sm, y_test)
print(f'Training F1 score: {f1_train}, Testing F1 score: {f1_test}')
```

	precision	recall	f1-score	support
0.0	0.97	0.90	0.93	95439
1.0	0.68	0.88	0.77	23370

accuracy			0.89	118809
macro avg	0.82	0.89	0.85	118809
weighted avg	0.91	0.89	0.90	118809

Training F1 score: 0.8950805398345668, Testing F1 score: 0.7658241921138453

We can see that training F1-score after SMOTE has significantly increased from 0.779 to 0.895 and the testing score has decreased from 0.779 to 0.765.


```
def precision_recall_curve_plot(y_test, pred_proba_c1):
    precisions, recalls, thresholds = precision_recall_curve(y_test,
pred_proba_c1)

    threshold_boundary = thresholds.shape[0]
    # plot precision
    plt.plot(thresholds, precisions[0:threshold_boundary],
linestyle='--', label='precision')
    # plot recall
    plt.plot(thresholds, recalls[0:threshold_boundary],
label='recalls')

    start, end = plt.xlim()
    plt.xticks(np.round(np.arange(start, end, 0.1), 2))

    plt.xlabel('Threshold Value'); plt.ylabel('Precision and Recall
Value')
    plt.legend(); plt.grid()
    plt.show()

precision_recall_curve_plot(y_test, model_smote.predict_proba(X_test)
[:,1])
```



Threshold value has increased after SMOTE from 0.35 to 0.7

Hyperparameter Tuning

L2

```
model_l2 = LogisticRegression(C = 5, penalty = 'l2', solver =
'liblinear')
f1_train, f1_test = training(model_l2, X_sm, X_test, y_sm, y_test)
print(f'Training F1 score: {f1_train}, Testing F1 score: {f1_test}')
```

	precision	recall	f1-score	support
0.0	0.97	0.90	0.93	95439
1.0	0.68	0.88	0.77	23370
accuracy			0.89	118809
macro avg	0.82	0.89	0.85	118809
weighted avg	0.91	0.89	0.90	118809

Training F1 score: 0.8951170091260161, Testing F1 score: 0.7657390595471893

Tuning lambda values (C) with penalty L2

```
lambda_values = [0.1, 0.01, 0.001, 0.0001, 10, 100]
```

```
for lambda_val in lambda_values:  
    model_l2 = LogisticRegression(C = lambda_val, penalty = 'l2', solver  
    = 'liblinear')  
    f1_train, f1_test = training(model_l2, X_sm, X_test, y_sm, y_test)  
    print('Lambda value: ', lambda_val)  
    print(f'Training F1 score: {f1_train}, Testing F1 score: {f1_test}')
```

	precision	recall	f1-score	support
0.0	0.97	0.90	0.93	95439
1.0	0.68	0.88	0.77	23370
accuracy			0.89	118809
macro avg	0.82	0.89	0.85	118809
weighted avg	0.91	0.89	0.90	118809

Lambda value: 0.1

Training F1 score: 0.8950678523812303, Testing F1 score:
0.766201987246033

	precision	recall	f1-score	support
0.0	0.97	0.90	0.93	95439
1.0	0.68	0.88	0.77	23370
accuracy			0.89	118809
macro avg	0.82	0.89	0.85	118809
weighted avg	0.91	0.89	0.90	118809

Lambda value: 0.01

Training F1 score: 0.8946107407590543, Testing F1 score:
0.7670606601248885

	precision	recall	f1-score	support
0.0	0.97	0.90	0.93	95439
1.0	0.68	0.87	0.77	23370
accuracy			0.89	118809
macro avg	0.82	0.89	0.85	118809
weighted avg	0.91	0.89	0.90	118809

Lambda value: 0.001

Training F1 score: 0.8913456563342349, Testing F1 score:
0.7658952496954933

	precision	recall	f1-score	support
0.0	0.97	0.89	0.93	95439
1.0	0.66	0.87	0.75	23370

accuracy			0.89	118809
macro avg	0.81	0.88	0.84	118809
weighted avg	0.90	0.89	0.89	118809

Lambda value: 0.0001

Training F1 score: 0.8833002984427699, Testing F1 score: 0.7491695578356832

	precision	recall	f1-score	support
0.0	0.97	0.90	0.93	95439
1.0	0.68	0.88	0.77	23370

accuracy			0.89	118809
macro avg	0.82	0.89	0.85	118809
weighted avg	0.91	0.89	0.90	118809

Lambda value: 10

Training F1 score: 0.895126329488848, Testing F1 score: 0.7657390595471893

	precision	recall	f1-score	support
0.0	0.97	0.90	0.93	95439
1.0	0.68	0.88	0.77	23370

accuracy			0.89	118809
macro avg	0.82	0.89	0.85	118809
weighted avg	0.91	0.89	0.90	118809

Lambda value: 100

Training F1 score: 0.8951099163706618, Testing F1 score: 0.7657390595471893

The F1-scores for train and test are almost same for all the values though highest for 0.1

L1

```
model_l1 = LogisticRegression(C = 5, penalty = 'l1', solver = 'liblinear')
```

```
f1_train, f1_test = training(model_l1, X_sm, X_test, y_sm, y_test)
```

```
print(f'Training F1 score: {f1_train}, Testing F1 score: {f1_test}')
```

	precision	recall	f1-score	support
0.0	0.97	0.90	0.93	95439
1.0	0.68	0.88	0.77	23370

accuracy			0.89	118809
macro avg	0.82	0.89	0.85	118809
weighted avg	0.91	0.89	0.90	118809

```
Training F1 score: 0.8951174798216542, Testing F1 score:
0.7658076588176446
```

L1 and L2 both are giving almost same F1-scores for train and test data.

Class weight

Using class weight algorithm to use the weight calculated as (no of samples in majority class / no of samples in minority class) for minority class and 1 for majority class

```
wt = y_train.value_counts()[0]/y_train.value_counts()[1]
wt
4.105077067565328
model = LogisticRegression(class_weight={0: 1, 1: wt})
model.fit(X_train, y_train)
LogisticRegression(class_weight={0: 1, 1: 4.105077067565328})
pred_y_test = model.predict(X_test)
```

Testing F1-score

```
f1_score(y_test, pred_y_test)
0.7629366641526094
```

Training F1-score

```
f1_score(y_train, pred_y_train)
0.7793166011490779
```

Both train and test F1-scores are poor as compared to other options that we tried

Thus, we can see that simple logistic regression model is giving higher F1-score of 0.779 for testing but SMOTE has significantly better training F1-score of ~0.9 and testing F1-score of 0.766

#Actionable insights and recommendations

- Around 80% of customers have fully paid their Loan amount. The defaulters are ~20%. From Personal loan business perspective, this ratio is high. These 20% will contribute in NPAs of LoanTap. To reduce the risk of NPAs: ** LoanTap should add slightly stringent rules to bring down this ratio to 5% to 6%. ** LoanTap should provide loans at slightly higher rate than other banks. This will offset the risks of defaulters and maintain the profitability of the business.

- The loan term 60 months has negative coefficient which means more chances of unlikely to pay. Which means LoanTap should focus more on loans for shorter duration (i.e. 36 months). Their social media campaign and marketing strategy should be based on this consideration.
- Overall statistics of the model [Classification Metrics]: ** Accuracy --> 92% ** Precision --> 86% ** Recall --> 71% ** F1-score --> 78%
- Precision is higher than recall which means that false positives is lesser than false negatives. That means LoanTap will not lose the potential clients that much with this model but might struggle in identifying NPAs.
- Features which have significant impact on outcome are as follow:
 1. int_rate: Interest Rate
 2. sub_grade: loan subgrade
 3. term : number of payments on the loan
 4. application_type
 5. zip code (from address)
 6. emp_title: job title supplied by the Borrower
- The sub_grade and grade logic to classify person by LoanTap is well created. From the model pov, it is considered to be significant. ** For the loan ratings 'A', 'B', 'C' and 'D', there's a huge difference between the no of defaulters and non-defaulters and no of non-defaulters is high which means more likely loan will be paid but for grades 'E', 'F' and 'G', these nos are almost same that means high risk ratings. ** Thus, we can say that grades 'A', 'B' and 'C' are low risk, grades 'D' and 'E' are moderate risk and 'F' and 'G' are high risk. ** Similar pattern observed for the sub-grades with 1 being low risk in that grade and 5 being high risk. ** So overall, A1 is lowest risk and G5 is the highest risk.

So it's recommended to avoid approving the loan for highest risk customers.

- Distribution of loan status across States: State codes 'AP', 'AE' and 'AA' are the top 3 states from which loan applications have been received (in same order). For all other states, it's almost similar. Thus, distribution is different across states. So, it's recommended that LoanTap should focus more on getting more customers from these 3 states as they are also more likely to repay the loan.
- Distribution of loan status across zip codes: ** We can see that the distribution of borrowers w.r.t. their loan status is significantly different as per the zip codes. Zip codes: 05113, 00813, 29597 are having only non-defaulters where as the zip codes: 11650, 86630, 93700 are having only defaulters. So, it's recommended that LoanTap should focus more on getting customers from zip codes 05113, 00813, 29597.
- Among the borrowers with initial list status as 'w' i.e. whole the difference between defaulters and non-defaulters is lesser as compared to 'f' which means that more chance of defaulting when entire amount is approved.

So, it's recommended to approve only fractional amount in the beginning.

- Employee tenure, income verification status, purpose of loan, home ownership status and loan application type do not make much impact on defaulting.

- Chances of defaulting is the least for home_ownership as 'Mortgage' so recommended to gather more borrowers having this criteria met
- We can see that when either negative records on borrower's public credit profile are present or bankruptcy records are available for borrower, then there are more chances of defaulting.

So, it's recommended to avoid approving the loan when negative records are found or bankruptcy is found.

```
'''
Questionnaire:

1. What percentage of customers have fully paid their Loan Amount?
Ans--> We can see that 80.4% are non-defaulters and 19.6% are defaulters.

2. Comment about the correlation between Loan Amount and Installment features.
Ans--> The correlation coefficient is the highest for loan amount and installment features which is +0.97 which tells us that greater the loan amount, greater will be the monthly installment amount.

3. The majority of people have home ownership as _____.
Ans--> 'Mortgage'

4. People with grades 'A' are more likely to fully pay their loan. (T/F)
Ans--> True

5. Name the top 2 afforded job titles.
Ans--> 'Teacher' followed by 'Manager'

6. Thinking from a bank's perspective, which metric should our primary focus be on..
ROC AUC
Precision
Recall
F1 Score
Ans--> The best metric to consider is : F1-score
As we need to give equal importance to both precision and recall.
We don't want to miss potential customers and at the same time we also don't want to give loan to defaulters

7. How does the gap in precision and recall affect the bank?
Ans-->
Precision 86% and Recall 71%
Precision is higher than recall which means that false positives is lesser than false negatives.
```


That means LoanTap will not lose the potential clients that much with this model might struggle in identifying NPAs.

From confusion matrix, % of misclassified points in each class:

class 0: 2.88%

class 1: 28.69%

If Recall value is low (i.e. FN is high), it means Bank's NPA (defaulters) may increase.

If Precision value is low (i.e. FP is high), it means Bank is losing in opportunity cost.

8. Which were the features that heavily affected the outcome?

Ans--> Features with higher positive coefficients:

zip_code

emp_title

title

term

sub_grade

revol_util

int_rate

open_acc

dti

home_ownership

loan_amnt

9. Will the results be affected by geographical location? (Yes/No)

Ans--> Yes

Distribution across States:

State codes 'AP', 'AE' and 'AA' are the top 3 states from which loan applications have been received (in same order).

For all other states, it's almost similar. Thus, distribution is different across states.

Distribution across zip codes:

We can see that the distribution of borrowers w.r.t. their loan status is significantly different as per the zip codes.

Zip codes: 05113, 00813, 29597 are having only non-defaulters where as the zip codes: 11650, 86630, 93700 are having only defaulters.

...